Data Platform Overview

motivation

The Data Platform provides tools for managing the data captured in an experimental research facility, such as a particle accelerator. The data are used within control systems and analytics applications, and facilitate the creation of machine learning models for those applications.

The Data Platform is agnostic to the source and acquisition of the data. A project goal is to manage data captured from the EPICS "Experimental Physics and Industrial Control System", however, use of EPICS is not required. The Data Platform APIs are generic and can be used from essentially all programming languages and any type of application.

requirements and objectives

- Provide an API for ingestion of heterogeneous data including scalar values, arrays, structures, and images.
- Handle the data rates expected for an experimental research facility such as a particle accelerator. A baseline performance requirement is to handle 4,000 scalar data sources sampled at 1 KHz, or 4 million samples per second.
- Provide an API for retrieval of ingested data.
- Provide an API for exploring metadata for data sources available in the archive.
- Provide mechanisms for adding post-ingestion annotations to the archive, and performing queries over those annotations.
- Provide mechanism for exporting data from the archive to common formats.

data platform elements

The Data Platform includes two primary technical components, including an API utilizing the gRPC framework, and a suite of services built using the Java programming language. The project also includes utilities for deploying and managing those components. A JavaScript web application for exploring the data archive and working with data is under development. Each element is described in more detail below.

gRPC API

The Data Platform API is built upon the gRPC open-source high-performance remote procedure call (RPC) framework. As described on Wikipedia, "this framework was originally developed by Google for use in connecting microservices. It uses HTTP/2 for transport, protocol buffers as the interface description languages, and provides features such as authentication and bidirectional streaming. It generates cross-platform client and server bindings for many languages."

We chose to use the gRPC framework for the Data Platform API because it can meet our performance requirements for data ingestion, and bindings are provided for virtually any programming language.

The API definition is managed separately from the service implementations so that it can be utilized for building client applications that are independent of other Data Platform technology. The Data Platform gRPC API is described in more detail in section "Data Platform API".

service implementations

The Data Platform Services are implemented as Java server applications. There are three independent server applications, providing ingestion, query, and annotation services, respectively. The MongoDB document-oriented database management system is used by the services for persistence. Section "Data Platform Service Implementations" provides more detail about the Java service implementations and the frameworks used to build them.

web application

The Data Platform Web Application is under development using the JavaScript React library. It will provide a user interface for navigating archive metadata and time-series data, viewing and creating annotations, and other tools for visualizing and exporting data.

installation and deployment support tools

A set of utilities is provided to help manage the Data Platform ecosystem. There are scripts for managing infrastructure services including MongoDB and the Envoy proxy (used for deploying the web application), and a set of simple process-management utilities for managing the Data Platform server and benchmark applications.

status and milestones

"datastore" prototype (2022)

A prototype implementation of the Data Platform services was built focusing on the creation of a general API supporting ingestion and query of heterogeneous data types including scalar, array / table, structure, and image. Service implementations were created using Java for both the Ingestion and Query Services, as well as libraries for building client applications. The prototype technology stack included both InfluxDB (for time series data) and MongoDB (for metadata). This prototype successfully demonstrated the use of gRPC APIs for ingestion and retrieval of heterogeneous, but did not meet the baseline performance requirements.

datastore web application prototype (2022)

The datastore prototype included development of a web application using JavaScript React and Tailwind libraries. The prototype provided simple user interfaces for navigating metadata, as well as querying and displaying time-series data. It demonstrated calling gRPC APIs from a browser-based application using the gRPC Web JavaScript implementation of gRPC for browser clients.

technology performance benchmarking (September 2023)

Performance benchmark applications were developed and utilized to evaluate candidate technologies for use in the Data Platform implementation in light of the project performance goal stated above. Benchmarks focused on gRPC for API communication; InfluxDB, MongoDB and MariaDB for database storage; and writing JSON and HDF5 files to disk.

The benchmark results showed that it was likely we could build service implementations meeting our performance requirements by using gRPC for communication and MongoDB for storing "buckets" of time series data.

Data Platform v1.0 (November 2023)

Version 1.0 of the Data Platform includes an initial Java implementation of the Ingestion Service providing a gRPC API and using MongoDB for storing time-series data. The initial ingestion service implementation focuses only on scalar data and with timestamps specified using the "sampling clock" mechanism with start time and sample period. It is accompanied by a performance benchmark application that is used at each stage of development to measure ingestion performance relative to the project goal. The initial implementation exceeds our goal by a comfortable margin, but this will continue to be a focus as the project evolves. section "Data Platform API" provides more information about the ingestion API.

v1.1 (January 2024)

Version 1.1 includes a Java implementation of the Query Service gRPC API, using the MongoDB database managed by the ingestion service to fulfill client query requests. A variety of API RPC methods for querying time-series data are provided to support the development of clients with varying performance requirements, ranging from streaming methods that return bucketed result data down to simple single response methods that return tabular data. See section "Data Platform API" for a detailed description of the query API.

v1.2 (February 2024)

Version 1.2 saw changes to the "proto" files defining the gRPC API for the Data Platform to be more consistent and conventional, with corresponding changes to the Java service implementations.

v1.3 (April 2024)

Version 1.3 provides an initial implementation of the annotation service for adding annotations to archived data and performing queries against those annotations. The primary focus for the initial annotation service implementation was on the data model for associating annotations with data in the archive. The only type of annotation currently supported is a simple user comment, but we will be adding many other types of annotations using the same underlying data model. See section "Data Platform API" for more details about the annotation data model.

v1.4 (July 2024)

Version 1.4 adds Ingestion and Query Service support for all data types defined in the Data Platform API including scalars, multi-dimensional arrays, structures, and images. Support is also added to both services for ingesting and querying data with an explicit list of data timestamps to complement the existing support for specifying data timestamps using a SamplingClock (with start time, sample period, and number of samples). Both features utilize serialization of the protobuf DataColumn and DataTimestamps API objects as byte array fields of the MongoDB BucketDocument. This change improves ingestion performance significantly, while also reducing the MongoDB storage footprint and simplifying the codebase.

v1.5 (August 2024)

With version 1.5, we have now completed the implementation of the initial Data Platform API we defined at the outset for the Core Services. This version focuses on adding the remaining unimplemented Ingestion Service features including: unidirectional client-side streaming data ingestion API, API for registering providers, API for querying ingestion request status details, testing for handling of value status information, validation of data ingestion providers, as well as improvements to the performance benchmark framework. The java dp-grpc and dp-service projects are updated to use Java 21 and the latest versions of 3rd party libraries.

v1.6 (October 2024)

Version 1.6 includes a new Annotation Service API method for exporting time-series data from the archive to common file formats including HDF5, CSV, and XLSX (Excel). It also provides an enhancement to the annotations query API method for filtering annotations by dataset id, in addition to the previously supported methods for filtering by owner and comment text field content.

todo and road map

- Add support for additional annotation types, including linked and derived data sets, and post-ingestion calculations.
- Run more extensive load testing benchmarks.
- Build simple data generator for demo and web application development.
- Implement mechanism for ingestion data validation.
- Add API for time-series data query by value status information?
- Add API for provider metadata query.
- Add framework for measuring data statistics.
- Add support for authentication and authorization of query and annotation services.
- Investigate MongoDB database clustering (replica sets), partitioning (sharding), and connection pooling.
- Experiment with horizontal scaling alternatives.
- Experiment with streaming architecture (e.g., Apache Kafka)

project organization

The Data Platform project is organized using the following github repositories:

dp-grpc

The dp-grpc repo contains the Data Platform API definition. It includes documentation for the Platform's data and service models, and a description of the gRPC "proto" files containing the API definition, which is provided in section "Data Platform API" of this document.

dp-service

The dp-service repo contains the Java code for implementations of the Data Platform services, including the shared frameworks used to build them. It includes documentation about those frameworks and the underlying MongoDB database schema utilized by the services, which is provided by Section "Data Platform Service Implementations" of this document.

dp-web-app

The dp-web-app repo contains the JavaScript code for the Data Platform Web Application, with documentation about the approach.

dp-support

The dp-support repo contains the scripts and utilities for managing the components of the Data Platform ecosystem. It includes documentation for using those tools.

data-platform

The data-platform repo is the primary repo for the Data Platform project. It contains documentation about the approach with links to the other repos, provided in Section "project organization" of this document. It also includes a Quick Start guide for running the Data Platform ecosystem from the installer.

dp-benchmark

The dp-benchmark repo is not currently active, but contains code developed for evaluating the performance of some candidate technologies considered for use in the Data Platform service technology stack. It includes an overview of the benchmark process with a summary of results.

Data Platform API

This section provides additional background for the gRPC framework used to implement the Data Platform API, the data and service models reflected in the API, and mapping of the API elements to the gRPC "proto" files that define the API.

gRPC background

gRPC is a framework that allows a client application to call a method on a server application. Defining an API with gRPC consists of identifying the services to be provided by the application, specifying the methods that can be called remotely for each service along with the method parameters and return types.

Underlying the gRPC framework is another Google-developed technology, Protocol Buffers, which is an open source mechanism for serializing structured data. gRPC uses Protocol Buffers as both the Interface Definition Language (IDL), and as the underlying message interchange format.

The gRPC API is defined using "proto" files (a text file with a ".proto" extension). Proto files contain definitions for services, service methods, and the data types used by those methods. Data types are called "messages", and each message specifies a series of name-value pairs called "fields". The definition of one message can be nested within another, limiting the scope of the nested data type to the message it is nested within.

See the links above for some simple examples of services, methods, and messages.

Data Platform gRPC API proto files

Currently, the Data Platform API defines three application services: ingestion, query, and annotation. The methods and data types (messages) for each service are contained in individual "proto" files (e.g., "ingestion.proto", "query.proto", and "annotation.proto"), with some shared data types in "common.proto" that are included in the relevant service files via "import" statements.

Data Platform proto file conventions

ordering of elements

Within the Data Platform service proto files, elements are listed in the following order:

1. service method definitions
2. definition of request and response data types
3. definition of other shared data types

packaging of parameters for a method into a single "request" message

For all Data Platform service methods, parameters are bundled into a single "request" message data type, instead of listing multiple parameters to the method.

naming of request and response messages

The service-specific proto files each begin with a "service" definition block that defines the method interface for that service, including parameters and return types. Where possible, the data types for the request and response use message names based on the corresponding method name.

A simple example is the Ingestion Service method "registerProvider()". The method request parameters are bundled in a message data structure called "RegisterProviderRequest". The method returns the response message type "RegisterProviderResponse". So the method definition looks like this:

```
rpc registerProvider (RegisterProviderRequest) returns (RegisterProviderResponse);
```

A more complex example is the Ingestion Service RPC methods "ingestDataStream()" (bidirectional streaming data ingestion API) and "ingestData()" (unary data ingestion API). We want both methods to use the same request and response data types, so we use the common message types "IngestionRequest" and "IngestionResponse". This pattern is also used for time-series data queries defined in "query.proto". The method definitions look like this:

```
rpc ingestData (IngestDataRequest) returns (IngestDataResponse);
rpc ingestDataStream (stream IngestDataRequest) returns (stream IngestDataResponse);
```

nesting of messages

Where possible, nesting is used to enclose simpler messages within the more complex messages that use them. In cases where we want to share messages between multiple request or response messages, the definition of those messages appears after the request and response messages in the proto file.

determining successful method execution

A common pattern is used across all Data Platform service method responses to assist in determining whether an operation succeeded or failed. All response messages use the gRPC "oneof" mechanism so that the message payload is either an "ExceptionalResult" message indicating that the operation failed, or a method-specific message containing the result of a successful operation.

The "ExceptionalResult" message is defined in "common.proto" with an enum indicating the status of the operation and a descriptive message. The enum indicates operations that were rejected, encountered an error in processing, failed to return data, resources that were unavailable when requested, etc.

Here is an example of the use of this pattern in the "QueryDataResponse" message used to send the result of time-series data queries:

```
message QueryDataResponse {

  oneof result {
    ExceptionalResult exceptionalResult = 10;
    QueryData queryData = 11;
  }

  message QueryData {

    repeated DataBucket dataBuckets = 1;

    message DataBucket {
      // DataBucket field definitions...
    }
  }
}
```

Data Platform API data model

The purpose of this section is to introduce some of the elements of the Data Platform's data model. These concepts will be used in subsequent descriptions of the various service APIs.

process variables

The core element of the Data Platform is the "process variable" (PV). In control theory, a process variable is the current measured value of a particular part of a process that is being monitored or controlled. The primary purpose of the Data Platform Ingestion and Query Services is to store and retrieve PV measurements. It is assumed that each PV for a particular facility is uniquely named. E.g., "S01:GCC01" might identify the first vacuum cold cathode gauge in sector one in the storage ring for some accelerator facility.

data vectors

The Data Platform Ingestion and Query Service APIs for handling data work with vectors of PV measurements. In "common.proto", this is reflected in the message data type "DataColumn", which includes a PV name and list of measurements.

handling heterogeneous data

One requirement for the Data Platform API is to provide a general mechanism for handling heterogeneous data types for PV measurements including simple scalar values, as well as multi-dimensional arrays, structures, and images. This is accomplished by the "DataValue" message data type in "common.proto", which uses the "oneof" mechanism to support a number of different data types for the values in a data vector (DataColumn).

timestamps

Time is represented in the Data Platform API using the "Timestamp" message defined in "common.proto". It contains two components for the number of seconds since the epoch, and nanoseconds. As a convenience, the message "TimestampList" is used to send a list of timestamps.

data providers

A data provider is an infrastructure component that uses the Data Platform Ingestion Service API to upload data to the archive. Before sending ingestion requests with data, the provider must be registered with the Ingestion Service.

ingestion data frame

The message "IngestionDataFrame", defined in "ingestion.proto", is the primary unit of ingestion in the Data Platform API. It contains the set of data to be ingested, using a list of "DataColumn" PV data vectors (described above). It uses the message "DataTimestamps", defined in "common.proto", to specify the timestamps for the data values in those vectors.

"DataTimestamps" provides two mechanisms for specifying the timestamps for the data values.

A "TimestampList" (described above) may be used to send an explicit list of "Timestamp" objects. It is assumed that each PV data vector "DataColumn" is the same size as the list of timestamps, so that there is a data value specified for each corresponding time value.

A second alternative is to use the "SamplingClock" message, defined in "common.proto". It uses three fields to specify the data timestamps, with a start time "Timestamp", the sample period in nanoseconds, and an integer count of the number of samples. The size of each data vector "DataColumn" in the "IngestionDataFrame" is expected to match the sample count.

bucketed time-series data

We use the "bucket pattern" as an optimization for handling time-series data in the Data Platform API for query results, as well as for storing a vector of PV measurement values in MongoDB. A "bucket" is a record that contains all the measurement values for a single PV for a specified time range.

This allows a list of values to be stored as a single unit in the database and returned in query results, as opposed to storing and returning individual data values and requiring that each record contains both a timestamp and data value

(which effectively triples the record size for scalar data). This leads to a more compact database, smaller gRPC messages to send query results, and improved overall performance.

A simple example of the bucket pattern follows (a slightly modified version of an example taken from the link above), demonstrating bucketing of temperature sensor data. The first snippet shows three measurements, with one record per measurement:

```
{
    sensor_id: 12345,
    timestamp: ISODate("2019-01-31T10:00:00.000Z"),
    temperature: 40
}

{
    sensor_id: 12345,
    timestamp: ISODate("2019-01-31T10:01:00.000Z"),
    temperature: 40
}

{
    sensor_id: 12345,
    timestamp: ISODate("2019-01-31T10:02:00.000Z"),
    temperature: 41
}
```

With bucketing, we save the overhead of the sensor_id and timestamp in each record:

```
{
    sensor_id: 12345,
    start_date: ISODate("2019-01-31T10:00:00.000Z"),
    sample_period_nanos: 1_000_000_000,
    count: 3
    measurements: [ 40, 40, 41 ]
}
```

Bucketing is used to send the results of time-series data queries. The message "QueryDataResponse" in "query.proto" contains the query result in "QueryData", which contains a list of "DataBucket" messages. Each "DataBucket" contains a vector of data in a "DataColumn" message for a single PV, along with time expressed using "DataTimestamps" (described above), with either an explicit list of timestamps for the bucket data values, or a SamplingClock with start time and sample period.

ingestion request status information

For performance reasons, data ingestion requests are handled asynchronously by the Ingestion Service. Each "IngestDataRequest" sent via a data ingestion API method is either acknowledged or rejected immediately. A request that is accepted for processing by the service may subsequently encounter an error during processing. For that reason, a "RequestStatus" record is created in the database for each ingestion request received by the service, indicating the disposition of that request (e.g., success, rejected, or error). The "queryRequestStatus()" API method is used to query request status details.

datasets

When designing the Data Platform's Annotation Service, we found we needed a mechanism for specifying a collection of data in the archive as the subject of an annotation. We decided to add the notion of a "dataset", consisting of a list of "data blocks", where each "data block" specifies a list of PV names and a time range.

If you think of the entire data archive as a giant spreadsheet, with a column for each PV name and a row for each measurement timestamp, a "data block" specifies some region within that spreadsheet, and a "dataset" contains a collection of those regions. This is illustrated in the figure below.
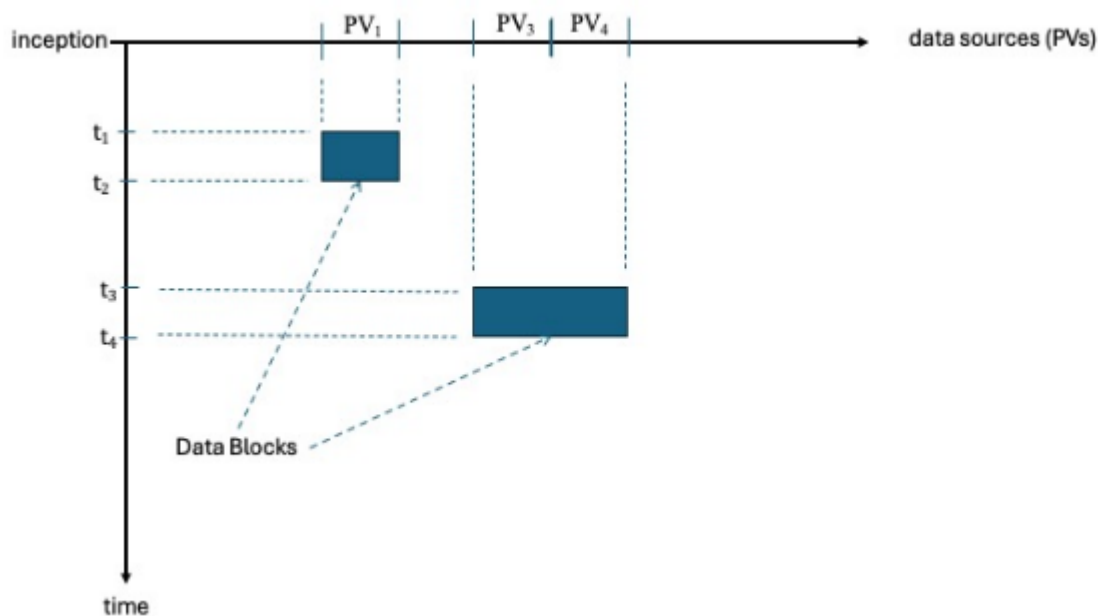
"annotation.proto" defines the messages "DataSet" and "DataBlock" for use as the data model for creating annotations, where a "DataSet" includes a description and a list of "DataBlock" messages, and each "DataBlock" includes begin and end Timestamp messages (described above), and a list of PV names.

annotations

An annotation allows clients to annotate the data archive with notes, data associations, and post-acquisition calculations. The Data Platform Annotation Service currently supports only a "comment" annotation, but additional types of annotations will be added in future releases.

Given the definition of a "DataSet" described above, the message "CreateAnnotationRequest" in "annotation.proto" is used to create an annotation, by providing the id of the "DataSet" to be annotated (which allows us to add multiple annotations to the same DataSet), and the details for the particular type of annotation to be created. In the case of a "CommentAnnotation", we simply specify the text of the comment.

For a link between related datasets, we might create a "LinkAnnotation" that specifies the id of the linked dataset and some text describing the relationship.

ingestion metadata

The Ingestion Service API allows descriptive metadata to be attached to data sent to the archive. Two types of metadata are supported, key/value attributes and event metadata.

The message "Attribute", defined in "common.proto", is a simple data structure that includes two strings, a key and a value. The message "IngestDataRequest", in "ingestion.proto" includes an optional list of "Attribute" messages that can be used to tag the request's data as it is added to the archive.

The message "EventMetadata", also defined in "common.proto", allows incoming data to be associated with some event. The "EventMetadata" message includes fields for the event description, with start and stop timestamps specifying the event start and stop time.

Data Platform API - ingestion service

"DpIngestionService" is a gRPC service defined in "ingestion.proto". It includes methods for provider registration and data ingestion.

provider registration

Data providers must be registered with the Ingestion Service before using ingestion API methods to send data to the ardhive. This is accomplished via the provider registration API:

```
rpc registerProvider (RegisterProviderRequest) returns (RegisterProviderResponse);
```

This unary method sends a single RegisterProviderRequest and receives a single RegisterProviderResponse. It is required to call this method to register a data provider before calling one of the data ingestion methods using the id of that provider.

Provider name is required in the RegisterProviderRequest, which may also contain optional metadata key/value attributes describing the provider.

The response message indicates whether the registration was successful. The response payload is an ExceptionalResult if the request is unsuccessful, otherwise it is a RegistrationResult that includes details about the new provider including providerId (for use in calls to data ingestion methods) and a flag indicating if the provider is new. On success, if a document already exists in the MongoDB "providers" collection for the provider name specified in the RegisterProviderRequest, the method returns the corresponding provider id in the response, otherwise a new document is created in the "providers" collection and its id returned in the response.

It is safe (and recommended) to call this method each time a data ingestion client is run. If a document already exists in the MongoDB providers collection for the specified provider, the attributes are updated to the values in the RegisterProviderRequest.

RegisterProviderRequest

Encapsulates parameters for a registerProvider() API method request. The required providerName field uniquely identifies a provider. The list of descriptive metadata key/value attributes is optional.

RegisterProviderResponse

Encapsulates response from registerProvider() API method. Payload is either an ExceptionalResult message with details about a rejection or error, or a RegistrationResult containing details about the provider registration including

provider id/name, and a flag indicating if the provider is new or previously existing.

The providerId contained in the RegistrationResult should be used for the corresponding parameter in IngestDataRequests sent to any of the data ingestion API methods. Those methods confirm that the specified providerId is valid before processing the ingestion request.

data ingestion

The Ingestion Service provides a streamlined API for ingesting data to the archive. There are three methods for data ingestion:

```
rpc ingestData (IngestDataRequest) returns (IngestDataResponse);
rpc ingestDataStream (stream IngestDataRequest) returns (IngestDataStreamResponse);
rpc ingestDataBidiStream (stream IngestDataRequest) returns (stream IngestDataResponse);
```

All methods use the same request message, "IngestDataRequest". The method "ingestData()" sends a single "IngestDataRequest" and receives a single "IngestDataResponse" corresponding to the request. "ingestDataStream()" is a unidirectional gRPC client-side streaming method that allows the client to send a stream of "IngestDataRequest" messages, and receive a single "IngestDataStreamResponse" message when the request stream is closed. "ingestDataBidiStream()" is a bidirectional gRPC streaming method that allows the client to send a stream of "IngestDataRequest" messages and receive a stream of "IngestDataResponse" messages. In all cases, the client uses the combination of provider id and request id to match incoming responses to outgoing requests.

ingestion data frame

As described above, the "IngestionDataFrame" is the primary unit of ingestion, containing a set of PV data vectors with the corresponding timestamp specification.

ingestion request

An "IngestDataRequest", defined in "ingestion.proto", includes an "IngestionDataFrame", optional metadata (list of key/value attributes or event metadata), a "Timestamp" indicating the time the request is sent, an id specifying the provider sending the data, and a mandatory client-generated request identifier, uniquely identifying the request for that provider.

ingestion response

The message "IngestDataResponse" in ingestion.proto contains one of two payloads, either an "ExceptionalResult" (described above) indicating an error or rejection, or an "AckResult", indicating the request was accepted and echoing back the dimensions of the request in confirmation. The response also includes provider id and client request id for matching the response to the corresponding request, and a "Timestamp" indicating the time the message was sent.

The Ingestion Service is fully asynchronous, so the response does not indicate if a request is successfully handled, only whether the request is accepted or rejected. The "queryRequestStatus()" API method is used to query request status information.

request status

For performance reasons, data ingestion requests are handled asynchronously by the Ingestion Service. Each "IngestDataRequest" sent via a data ingestion API method is either acknowledged or rejected immediately. A request that is accepted for processing by the service may subsequently encounter an error during processing. For that reason, a "RequestStatus" record is created in the database for each ingestion request received by the service, indicating the disposition of that request (e.g., success, rejected, or error). The "queryRequestStatus()" API method is used to query request status details:

```
rpc queryRequestStatus(QueryRequestStatusRequest) returns (QueryRequestStatusResponse);
```

This unary method sends a single QueryRequestStatusRequest and receives a single QueryRequestStatusResponse. It is used to determine the status of an individual data ingestion request, or to find data ingestion errors for a specified time range.

The QueryRequestStatusRequest message contains a list of criteria for searching by provider id, provider name, request id, status, and time range. The criteria can be combined arbitrarily, but we envision three primary use cases:

1. Query by provider id or name and request id to find the status of a specific ingestion request.
2. Query by provider id or name, status (e.g., rejected or error) and time range.
3. Query by status and time range without specifying a provider (e.g., "find all ingestion errors for today").

The QueryRequestStatusResponse message payload is either an ExceptionalResult containing details about a rejection or error, or a RequestStatusResult containing a list of RequestStatus messages, one for each document in the MongoDB "requestStatus" collection that matches the search criteria.

Each RequestStatus message contains details about the status of an individual ingestion request, including provider id/name, request id, status enum, status message, and list of bucket ids created (documents added to the MongoDB "buckets" collection).

QueryRequestStatusRequest

Encapsulates parameters to queryRequestStatus() API method. Includes a list of criteria that can be arbitrarily combined for the query filter.

- ProviderIdCriterion is used to filter by providerId (as returned by registerProvider()) and sent in the IngestDataRequest.
- ProviderNameCriterion is used to filter by providerName (as sent to registerProvider()) and sent in the IngestDataRequest.
- RequestIdCriterion is used to filter by clientRequestId as specified in the IngestDataRequest.
- StatusCriterion is used to filter by status indicated for the IngestDataRequest (success, rejected, or error).
- TimeRangeCriterion is used to filter by time range, matched against the time indicated for the IngestDataRequest.

QueryRequestStatusResponse

Encapsulates response from queryRequestStatus() API method. Payload is an ExceptionalResult if the query is rejected or there is an error handling it, otherwise payload is a RequestStatusResult which contains a list of RequestStatus messages, corresponding to the documents in the MongoDB "requestStatus" collection that match the specified query criteria.

Each RequestStatus message includes details from the document in the "requestStatus" collection, including a unique identifier for the request status document, provider id and name, client request id, status enum value, status message, and a list of idsCreated in the MongoDB "buckets" collection for the request's data.

Data Platform API - query service

"DpQueryService" is a gRPC service defined in "query.proto". It includes methods for querying both time-series data and metadata.

time-series data query

The Query Service provides several methods for querying time-series data, offering different options for performance and packaging of results.

```
rpc queryData(QueryDataRequest) returns (QueryDataResponse);
rpc queryDataTable(QueryDataRequest) returns (QueryTableResponse);
rpc queryDataStream(QueryDataRequest) returns (stream QueryDataResponse);
rpc queryDataBidiStream(stream QueryDataRequest) returns (stream QueryDataResponse);
```

Each method accepts a "QueryDataRequest" message, described in more detail below, to specify the query parameters. All the time-series data query methods return "QueryDataResponse" messages with the exception of queryDataTable(), which returns data in a tabular format via a "QueryTableResponse" message. Each of the time-series query methods and the corresponding request/response objects is discussed in more detail below.

queryData()

The "queryData()" method is a simple unary method with a single request and response. The result of the query must fit in a single gRPC response message, or an error is generated. The "QueryDataResponse" contains bucketed time-series data (described above).

queryDataTable()

Also a single request/response unary method, queryDataTable() differs in that the "QueryTableResponse" message contains a tabular result for use in clients such as the Data Platform Web Application that display data in tables.

queryDataStream()

Expected to be the best performing time-series data query method for retrieving a large amount of data, "queryDataStream()" accepts a single "QueryDataRequest" message and returns its result as a stream of "QueryDataResponse" messages, each of which contains bucketed time-series data.

queryDataBidiStream()

"queryDataBidiStream()" is a bidirectional streaming method. This method is similar to queryDataStream(), but is used in clients that need explicit control of the response stream due to memory or performance considerations.

The client sends a single "QueryDataRequest" message, receiving a single "QueryDataResponse" with bucketed time-series data. The client then requests the next response in the stream by sending a "QueryDataRequest" containing a

"CursorOperation" method with type set to "CURSOR_OP_NEXT" until the result is exhausted and the stream is closed by the service.

QueryDataRequest

All time-series data query methods accept a "QueryDataRequest" message (defined in "query.proto"). The message contains one of two payloads, either a "QuerySpec" or a "CursorOperation".

A "QuerySpec" message payload specifies the parameters for a time-series data query and includes begin and end timestamps specifying the time range for the query, and a list of PV names whose data to retrieve for the specified time range.

A "CursorOperation" payload is a special case and applies only to the "queryDataBidiStream()" method. It contains an enum value from "CursorOperationType" specifying the type of cursor operation to be executed. Currently, the enum contains a single option "CURSOR_OP_NEXT" which requests the next message in the response stream. We may add additional operations, e.g, "fetch the next N buckets".

QueryDataResponse

Except for "queryDataTable()", all time-series data query methods return "QueryDataResponse" messages. A "QueryDataResponse" contains one of two message payloads, "ExceptionalResult" if an error is encountered or no data is found (described above) or "QueryData".

A "QueryData" message includes a list of "DataBucket" messages. Each "DataBucket" contains a vector of data in a "DataColumn" message for a single PV, along with time expressed using "DataTimestamps" (described above), with either an explicit list of timestamps for the bucket data values, or a SamplingClock with start time and sample period. The "DataBucket" also includes a list of key/value "Attribute" messages and/or "EventMetadata" message if specified on the ingestion request that created the bucket.

QueryTableResponse

The "queryDataTable()" time-series data query method returns its result via a "QueryTableResponse" message. This is essentially a packaging of the bucketed time-series data managed by the archive into a tabular data structure for use in a client such as a web application. A "QueryTableResponse" object contains one of two payloads, an "ExceptionalResult" if an error is encountered or no data is found (described above) or a "TableResult".

A "TableResult" message contains a list of PV column data vectors, one for each PV specified in the "QueryDataRequest". It also contains a "DataTimestamps" message with a "TimestampList" of timestamps, one for each data row in the table. The column data vectors are the same size as the list of timestamps, and are padded with empty values where a column doesn't contain a value at the specified timestamp.

metadata query

The Data Platform Query Service includes a single method for querying the archive's metadata about the PVs available in the archive.

```
rpc queryMetadata(QueryMetadataRequest) returns (QueryMetadataResponse);
```

"queryMetadata()" is a single request/response unary method that accepts a "QueryMetadataRequest" and returns a "QueryMetadataResponse".

QueryMetadataRequest

The "QueryMetadataRequest" message is defined in query.proto, and contains one of two payloads, "PvNameList" or "PvNamePattern". A "PvNameList" message specifies an explicit list of PVs to find metadata for. A "PvNamePattern" specifies a regular expression pattern for matching against PV names available in the archive.

QueryMetadataResponse

The "QueryMetadataResponse" message contains the result of a metadata query and includes one of two payloads, either an "ExceptionalResult" if an error is encountered or no data is found (described above) or "MetadataResult".

A "MetadataResult" message contains a list of "PvInfo" messages, one for each PV specified for the query (either explicitly in the PV name list or by matching the supplied PV name pattern). A "PvInfo" message contains metadata for an individual PV in the archive, including name, timestamps for the first and last PV measurement in the archive, and stats for the most recent bucket including bucket id, data type information, data timestamps details, and sample count/period.

Data Platform API - annotation service

"DpAnnotationService" is a gRPC service defined in "annotation.proto". It includes methods for creating and querying "DataSets" (described above in , and for creating and querying annotations.

creating and querying datasets

The Data Platform Annotation Service uses datasets to identify the relevant data within the archive for a particular annotation. The API includes a methods for creating and querying datasets.

```
rpc createDataSet(CreateDataSetRequest) returns (CreateDataSetResponse);
rpc queryDataSets(QueryDataSetsRequest) returns (QueryDataSetsResponse);
```

createDataSet()

This is a single request/response unary method for creating a dataset. It accepts a "CreateDataSetRequest" message and returns a "CreateDataSetResponse".

CreateDataSetRequest

A "CreateDataSetRequest" message contains a "DataSet" message with details of the dataset to be created, e.g., its list of "DataBlock" messages.

CreateDataSetResponse

A "CreateDataSetResponse" message contains one of two payloads, an "ExceptionalResult" message if a rejection or error was encountered creating the dataset, or a "CreateDataSetResult".

A "CreateDataSetResult" message simply contains the unique identifier assigned to the new dataset if it was created successfully.

queryDataSets()

The "queryDataSets()" method is a single request/response method that searches for datasets in the archive that match the search criteria specified for the query. It accepts a "QueryDataSetsRequest" message and returns a "QueryDataSetsResponse" message.

QueryDataSetsRequest

A "QueryDataSetsRequest" encapsulates the criteria for the query. It contains a list of "QueryDataSetsCriterion" messages.

The "QueryDataSetsCriterion" message defines a number of different criteria message types that can be added to the criterion list, including an "OwnerCriterion" (specifying the owner id to match in the annotation query), "NameCriterion" (specifying text to match against dataset name), and "DescriptionCriterion" (specifying text to match against dataset description.

These query criteria can be used individually in the criteria list, or multiple criteria can be added to the list to specify a compound query. E.g., adding an "OwnerCriterion" and "NameCriterion" to the list will match dataset names for the specified owner.

QueryDataSetsResponse

The "queryDataSets()" method returns a "QueryDataSetsResponse" message with the query results. It contains one of two payloads, either an "ExceptionalResult" message if the query encountered an error or returned no data (described above), or an "DataSetsResult" message with the results if the query was successful.

The "DataSetsResult" message includes a list of "DataSet" messages, one for each dataset that matches the query's search criteria.

An "DataSet" message includes the following properties for the dataset: unique id, name, owner id, description, and a list of the "DataBlock" messages comprising the dataset.

creating and querying annotations

The Data Platform Annotation Service provides two methods related to annotations.

```
rpc createAnnotation(CreateAnnotationRequest) returns (CreateAnnotationResponse);
rpc queryAnnotations(QueryAnnotationsRequest) returns (QueryAnnotationsResponse);
```

createAnnotation()

The method "createAnnotation()" creates an annotation for the specified dataset. It accepts a "CreateAnnotationRequest" message and returns a "CreateAnnotationResponse" message.

CreateAnnotationRequest

A "CreateAnnotationRequest" message specifies the id of the owner creating the annotation, and the id of the dataset to be annotated. It uses a variable "oneof" payload for specifying the details specific to the type of annotation being created. Currently there is a single type of annotation, "CommentAnnotation", which includes the text of the comment for the annotation.

CreateAnnotationResponse

A "CreateAnnotationResponse" message is used to return the result of the "createAnnotation()" method. It includes one of two payloads, either an "ExceptionalResult" (described above) if an error is encountered creating the annotation, or a "CreateAnnotationResult" if the operation is successful.

A "CreateAnnotationResult" message simply contains the unique identifier assigned to the annotation.

queryAnnotations()

The "queryAnnotations()" method is a single request/response method that searches for annotations in the archive that match the search criteria specified for the query. It accepts a "QueryAnnotationsRequest" message and returns a "QueryAnnotationsResponse" message.

QueryAnnotationsRequest

A "QueryAnnotationsRequest" encapsulates the criteria for the query. It contains a list of "QueryAnnotationsCriterion" messages.

The "QueryAnnotationsCriterion" message defines a number of different criteria message types that can be added to the criterion list, including an "OwnerCriterion" (specifying the owner id to match in the annotation query), "DataSetCriterion" (specifying the dataset id to match in the annotation query), and "CommentCriterion" (specifying text to match against annotation comments). Other types of criterion messages will be added as additional types of annotations are defined.

These query criteria can be used individually in the criteria list, or multiple criteria can be added to the list to specify a compound query. E.g., adding an "OwnerCriterion" and "CommentCriterion" to the list will match comment annotations for the specified owner.

QueryAnnotationsResponse

The "queryAnnotations()" method returns a "QueryAnnotationsResponse" message with the query results. It contains one of two payloads, either an "ExceptionalResult" message if the query encountered an error or returned no data (described above), or an "AnnotationsResult" message with the results if the query was successful.

The "AnnotationsResult" message includes a list of "Annotation" messages, one for each annotation that matches the query's search criteria.

An "Annotation" message includes the unique id of the annotation, the owner id, the id of the associated dataset identifying the data in the archive that the annotation applies to, and for convenience (so that a second query to retrieve the dataset is not required) a "DataSet" message containing the list of "DataBlocks" comprising the annotation's dataset.

exporting datasets

The Data Platform Annotation Service provides a method for exporting data for a DataSet to common file formats.

```
rpc exportDataSet(ExportDataSetRequest) returns (ExportDataSetResponse);
```

exportDataSet()

The method "exportDataSet()" exports data for the specified dataset to a file. It accepts a "ExportDataSetRequest" message and returns a "ExportDataSetResponse" message.

ExportDataSetRequest

Parameters to the "exportDataSet()" method are encapsulated in an ExportDataSetRequest that includes fields for specifying the id of the DataSet to be exported, and an enum for specifying the desired output file format.

ExportDataSetResponse

The "exportDataSet()" method returns an ExportDataSetResponse, whose payload includes either an ExceptionalResult indicating a problem in handling the export request, or an ExportDataSetResult for a successful request.

ExportDataSetResult includes fields specifying the full path for the export output file and (optionally if configured) the URL for accessing the file via a web server.
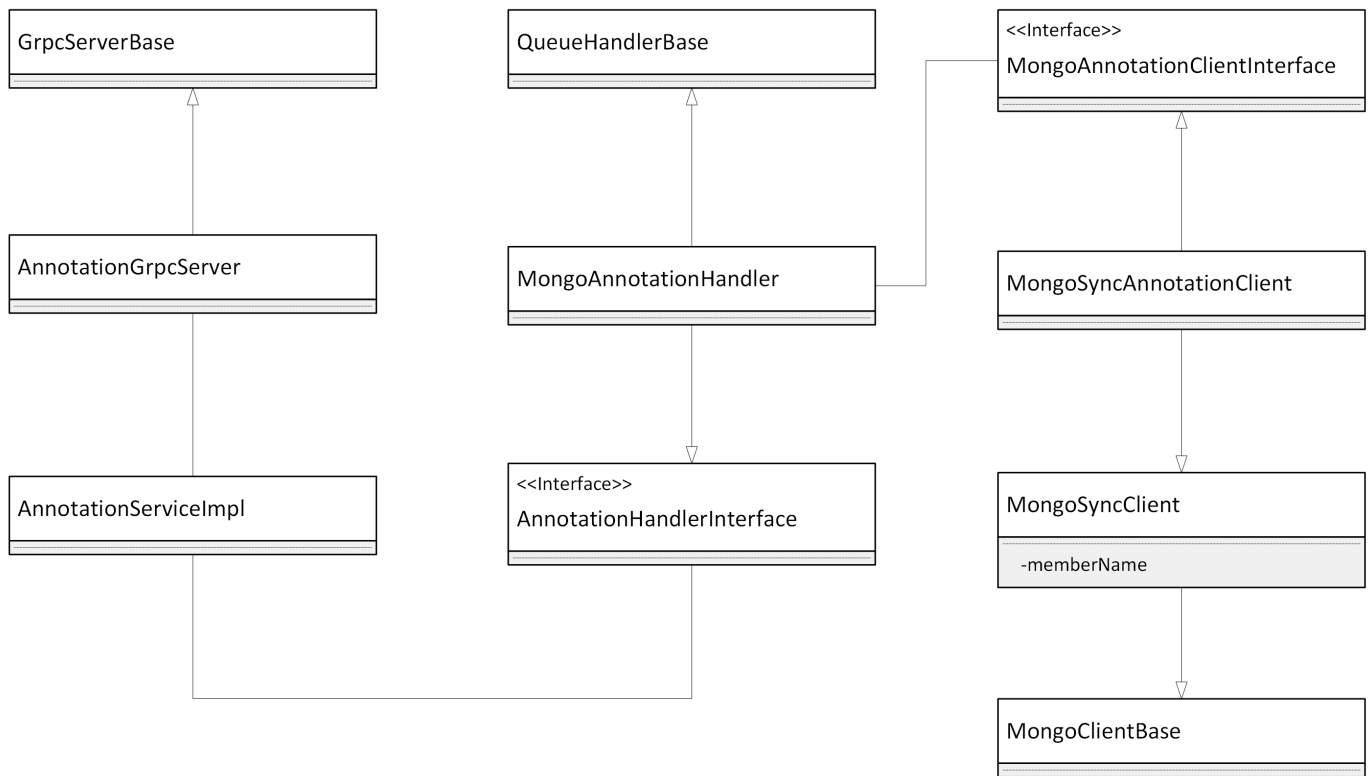
Data Platform Service Implementations

The dp-service repo contains Java implementations of the services defined in the dp-grpc repo's gRPC proto files, Ingestion, Query and Annotation. The remainder of this section covers the patterns and frameworks used to build the services (Section "dp-service patterns and frameworks") and some details about the MongoDB database schema (Section "dp-service MongoDB schema and data flow").

dp-service patterns and frameworks

Some common frameworks and patterns are used in the service implementations, including gRPC server, service request handling, MongoDB interface, configuration, performance benchmarking, regression testing, and integration testing.
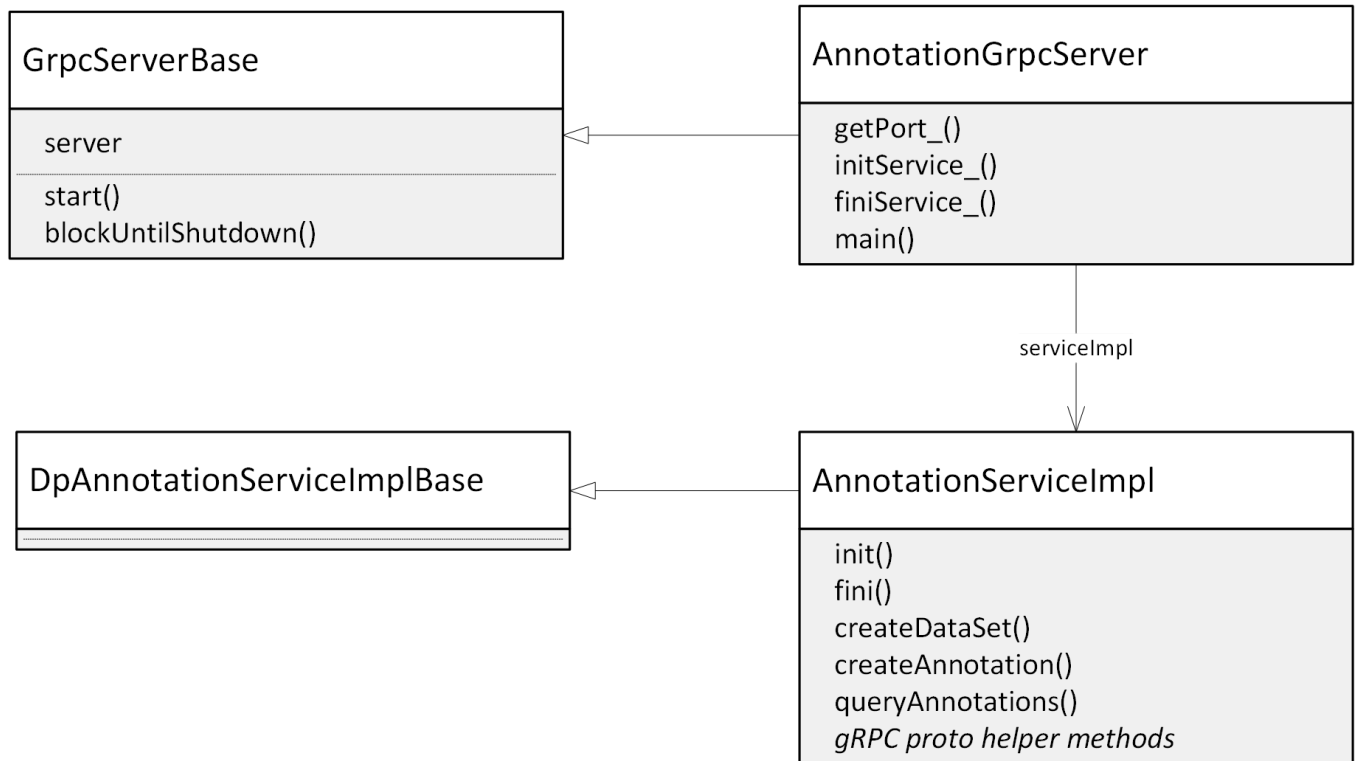
The diagram below gives an overview of the key classes used to build a service, using the concrete Annotation Service classes. The same pattern is used in all Data Platform service implementations, with analogous concrete classes for each.

GrpcServerBase		QueueHandlerBase		<<Interface>> MongoAnnotationClientInterface

AnnotationGrpcServer		MongoAnnotationHandler		MongoSyncAnnotationClient

AnnotationServiceImpl		<<Interface>> AnnotationHandlerInterface		MongoSyncClient -memberName

				MongoClientBase

The remainder of this section drills into each of the frameworks and patterns mentioned above, using the overview class diagram as a road map and exposing additional detail where appropriate.

gRPC server

Each service implementation includes an extension of the framework class "GrpcServerBase" that is the entry point for running the service. The diagram below shows the Annotation Service extension of that base class, "AnnotationGrpcServer". The extension implements the "main()" method to run the application.

GrpcServerBase

server

start()
blockUntilShutdown()

AnnotationGrpcServer

getPort_()
initService_()
finiService_()
main()

serviceImpl

DpAnnotationServiceImplBase

AnnotationServiceImpl

init()
fini()
createDataSet()
createAnnotation()
queryAnnotations()
gRPC proto helper methods

"GrpcServerBase" implements the method "start()" to initialize the gRPC communication framework, instantiate the class implementing the service implementation, and add a shutdown hook to handle VM shut down. It also implements "blockUntilShutdown()" to allow "main()" to wait for the application to complete and clean up the service.

The base class defines an abstract method interface to support the framework method implementations. "getPort_()" returns the port number for the service's gRPC server. initService_() and finiService_() are used to initialize and finalize the underlying service implementation.

For each service, the gRPC "protoc" compiler generates a service implementation stub class that defines the methods implemented by that service. The service extends that stub class to provide a useful implementation of the service API methods. In the case of the AnnotationService, this class is "AnnotationServiceImpl".

"AnnotationServiceImpl" provides methods "init()" and "fini()", for use by "initService_()" and "finiService_()", respectively, in "AnnotationGrpcServer" to initialize and finalize the service. It overrides methods that implement the gRPC service API, in this case "createDataSet()", "createAnnotation()", and "queryAnnotations()".

The handling for incoming requests by those methods is described in the next section.

service request handling framework

The diagram below shows the framework for handling incoming requests.
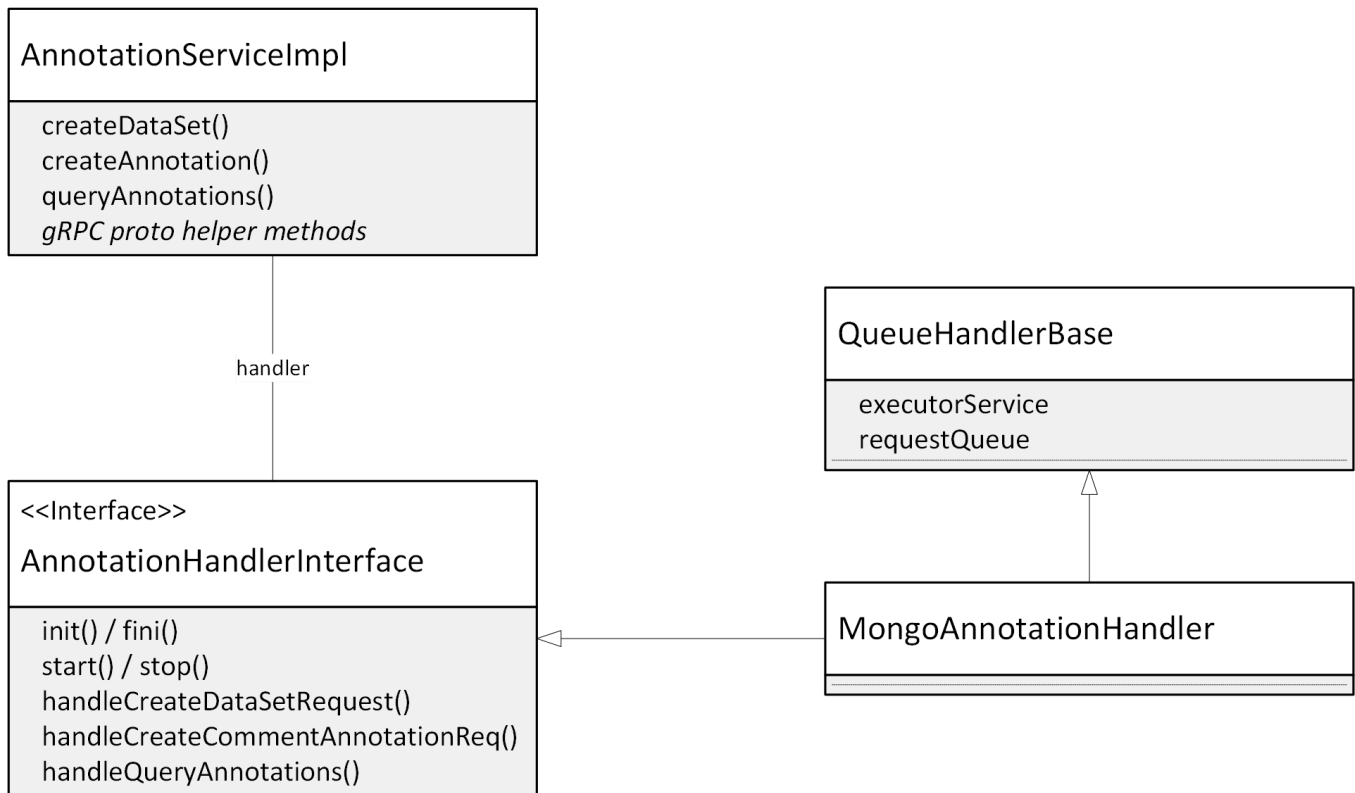
Requests coming in via the gRPC communication framework are delivered to the "AnnotationServiceImpl" for processing. Its main purpose is to provide useful implementations of the Annotation Service API methods. It also provides helper methods for the gRPC protocol related to the service, like creating response objects and sending them in the response stream.

Each service implementation defines a handler interface that defines the methods needed for handling requests. In the case of the AnnotationService, the interface is AnnotationHandlerInterface. It defines control methods for use by the "AnnotationServiceImpl" in managing the handler including "init()", "fini()", "start()", and "stop()". It defines methods for handling requests, including handleCreateDataSetRequest(), handleCreateCommentAnnotationReq(), and handleQueryAnnotations().

We use an interface to define the handler methods so that we don't prescribe a particular implementation. That said, our initial implementation uses MongoDB for persistence and the class "MongoAnnotationHandler" implements the "AnnotationHandlerInterface" by using MongoDB to store and retrieve annotations.

The framework provides the class "QueueHandlerBase" that provides a simple framework implementing the producer-consumer pattern that is extended in the concrete service handler implementation. It includes a queue for incoming tasks, and a pool of worker threads for processing them.

In addition to implementing the "AnnotationHandlerInterface" defining the service API methods, "MongoAnnotationHandler" also extends "QueueHandlerBase" to utilize the task queue and worker thread pool.

The next diagram provides additional detail for the "QueueHandlerBase" framework.
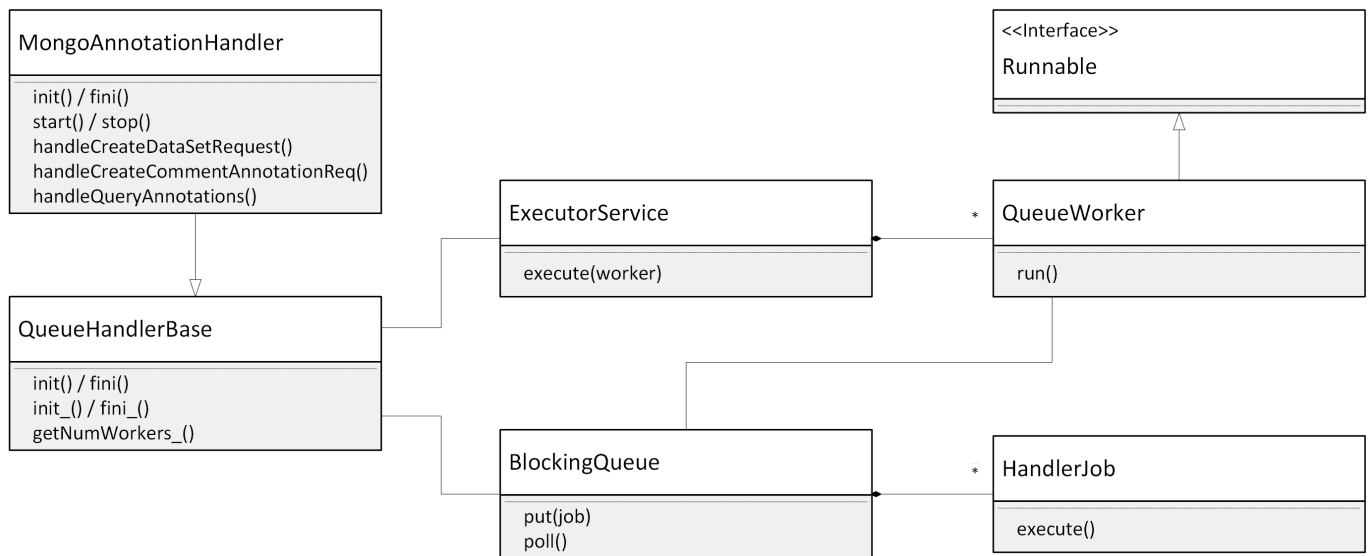
In "QueueHandlerBase" initialization, the concrete "MongoAnnotationHandler" implementation of abstract method "getNumWorkers_()" is used to create a Java "ExecutorService" containing the specified number of "QueueWorker" instances, starting each of them via "execute()". This invokes the workers' "run()" method, which monitors the handler's "BlockingQueue" for new tasks using poll().

The "MongoAnnotationHandler" methods for handling incoming requests, such as "handleQueryAnnotations()" create an instance of a service-specific subclass derived from HandlerJob, and add it to the queue using "put(job)". The next available "QueueWorker" removes the job from the queue, and processes it by calling the job's "execute()" method.

"HandlerJob" doesn't prescribe any particular implementation or relationship for derived concrete classes, and could probably be an interface instead of a class. Regardless, the intention is to support a variety of diffent types of jobs. The primary use case for a job is to 1) perform some database operation like creating an item or executing a query and 2) dispatch the results of that operation in the response stream for the API request.

The diagram below shows the static relationships for a "QueryAnnotationsJob", responsible for executing a "queryAnnotations()" API request and dispatching the results in the response stream.
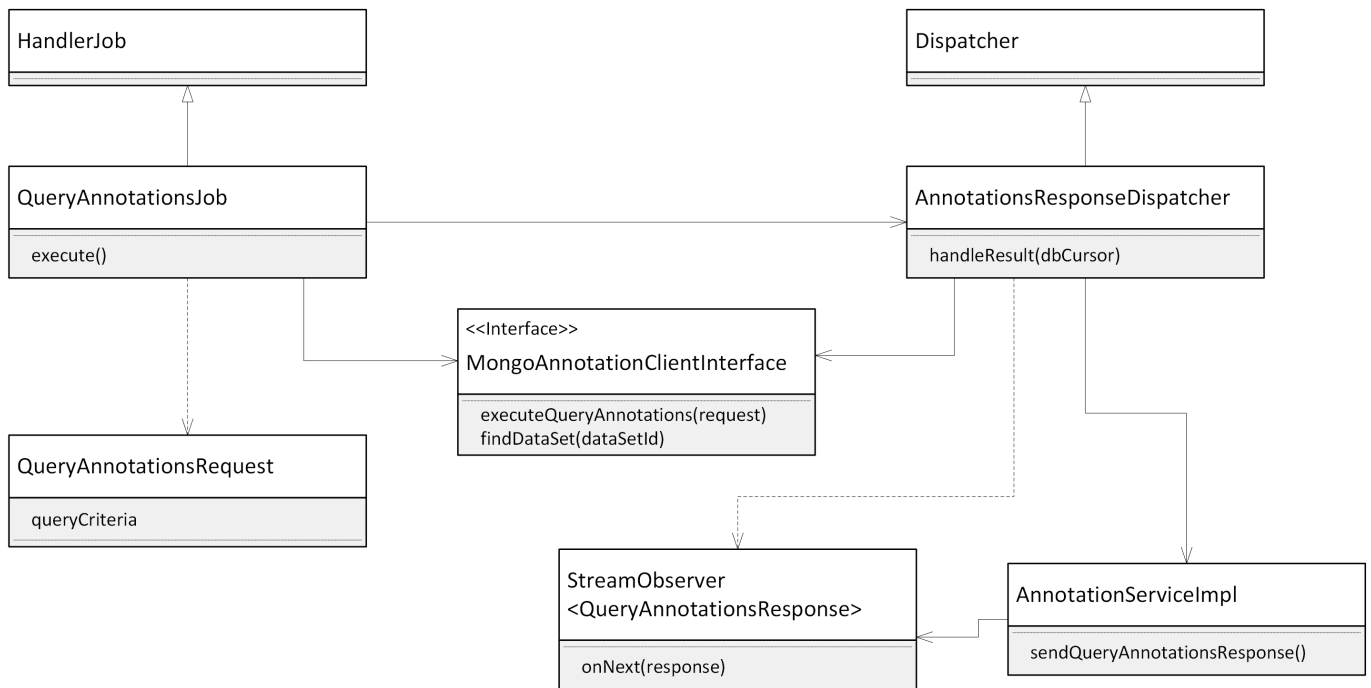
The "QueryAnnotationsJob" is removed from the handler's task queue by a worker, who invokes the job's "execute()" method. The job uses the handler's "MongoAnnotationClientInterface" to invoke "executeQueryAnnotations()" with the "QueryAnnotationsRequest", and passes the resulting database Cursor object to the "AnnotationsResponseDispatcher" method "handleResult()" for dispatching the query result.

The dispatcher uses a convenience method "sendQueryAnnotationsResponse()" on "AnnotationServiceImpl" to send the "QueryAnnotationsResponse" created by the dispatcher and containing the query results in the API response stream via the "StreamObserver" object's onNext() method.

This pattern for executing an API request and dispatching the response is used in all the service implementations for handling all requests. Instead of repeating diagrams like above for each case, they are summarized in the following tables listing service rpc method, handler method, job subclass, and dispatcher subclass, with a table for each service.

annotation service request handling

rpc method	handler method	job class	dispatcher
createDataSet	handleCreateDataSetRequest	CreateDataSetJob	CreateDataS
createAnnotation	handleCreateCommentAnnotationRequest	CreateCommentAnnotationJob	CreateAnno
queryAnnotations	handleQueryAnnotations	QueryAnnotationsJob	Annotationsl

ingestion service request handling

rpc method	handler method	job class	dispatcher class
ingestDataStream	onNext	TODO	TODO

query service request handling

| rpc method	handler method	job class	dispatcher class
queryData	handleQueryData	QueryDataJob	DataResponseUnaryDispatcher
queryDataStream	handleQueryDataStream	QueryDataJob	DataResponseStreamDispatcher
queryDataBidiStream	handleQueryDataBidiStream	QueryDataJob	DataResponseBidiStreamDispatch
queryTable	handleQueryTable	QueryTableJob	TableResponseDispatcher
queryMetadata	handleQueryMetadata	QueryMetadataJob	MetadataResponseDispatcher

handling for bidirectional streaming API methods

The core request handling pattern described in section "service request handling framework" is used to handle all requests, whether 1) unary with single request / response, 2) single request with server-side streaming response, or 3) bidirectional streaming with multiple requests and responses.

The third case is a special one, because we need a mechanism for directing subsequent requests from the API RPC method's request stream to the existing handler for the initial request. This is illustrated by the handling framework for the bidirectional time-series data query method, "queryDataBidiStream()".

That query method is a bidirectional streaming method, where the initial request specifies the time-series data query parameters, the initial response contains the first set of query results, and subsequent requests are sent to retrieve incremental results until the result set is exhausted.

Behind the scenes, a MongoDB query is executed to retrieve the relevant time-series data. The cursor for the query results is used to generate the initial result set in the response stream, but we want to hold on to the cursor to fulfill the subsequent requests for additional query results so that the query doesn't need to be re-executed. The diagram below illustrates the framework for handling "queryDataBidiStream()".
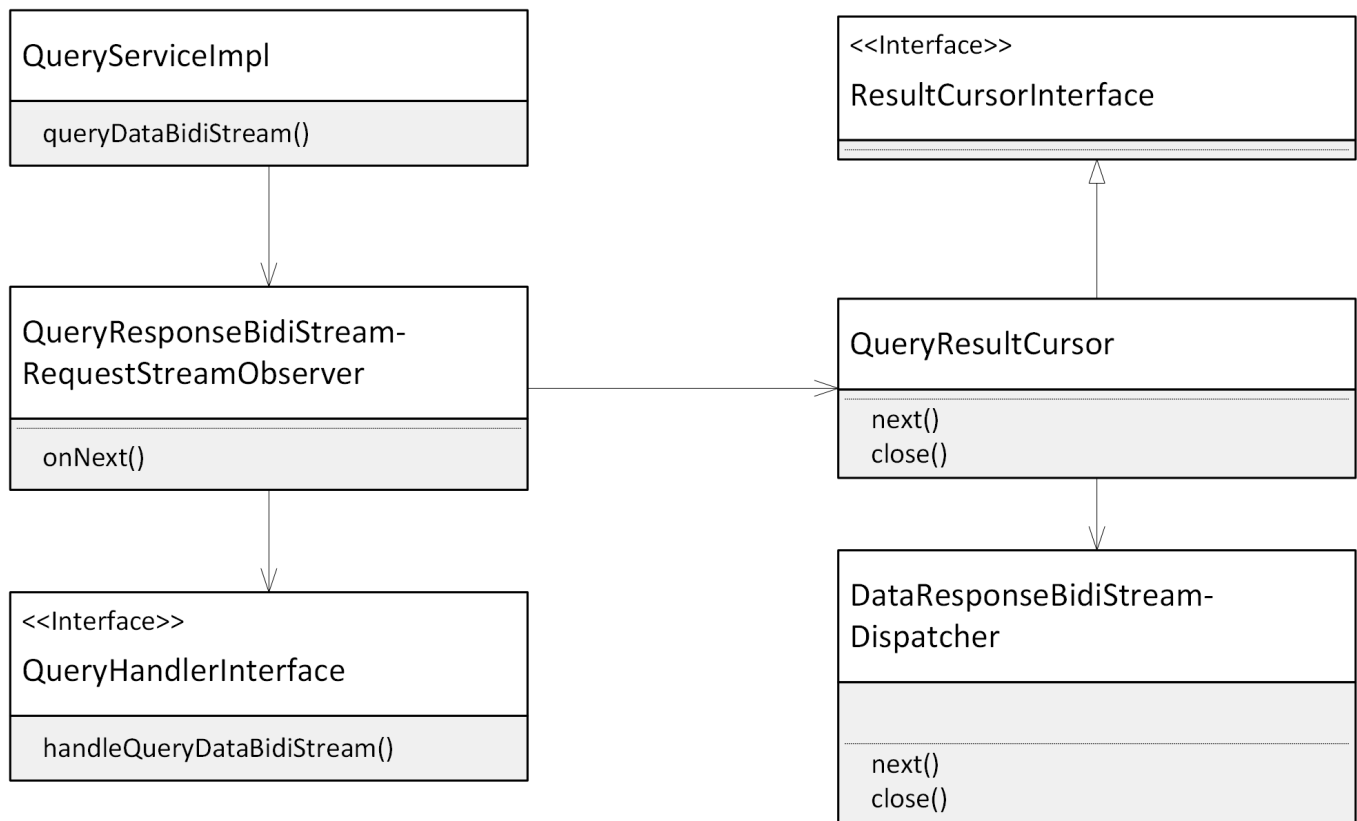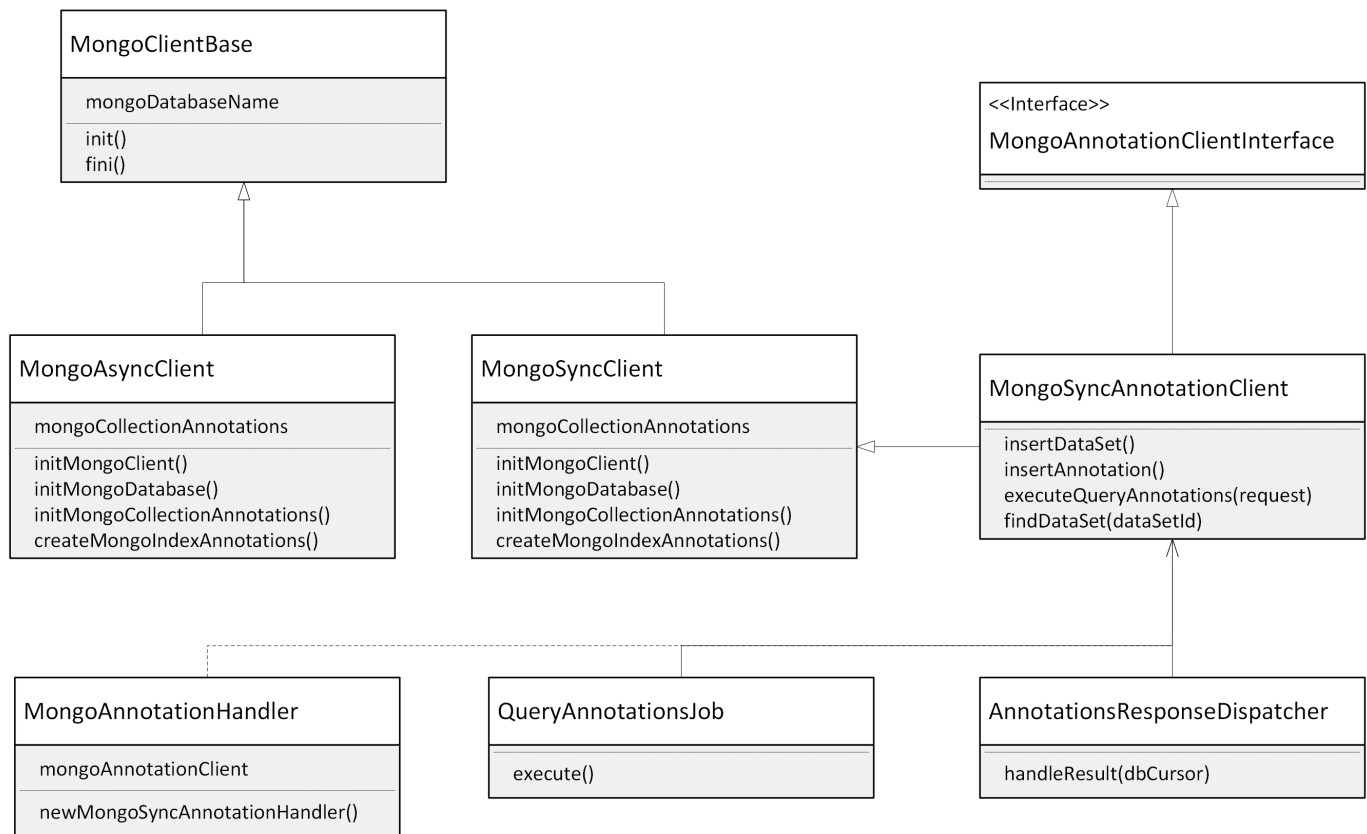
When "QueryServiceImpl.queryDataBidiStream()" receives a new request, it creates a
"QueryResponseBidiStreamRequestStreamObserver" to handle the request stream. The initial request received by its
"onNext()" method contains the parameters for the time-series data query and is dispatched to
"QueryHandlerInterface.handleQueryBidiStream()". This request is handled as illustrated in section "service request
handling framework". That method also returns a "QueryResultCursor" to the request stream observer.

The initial query results are returned in the response stream by the dispatcher. Subsequent requests arriving in the
request stream observer's "onNext()" method contain a cursor operation to retrieve the next set of results and are
dispatched to "QueryResultCursor.next()", which creates the next response message and sends it via the dispatcher.
This continues until the result set is exhausted and / or the API stream is closed.

MongoDB interface

MongoDB is used at the core of all Data Platform service implementations for data persistence. section "service request
handling framework" showed how the handler framework for the annotation service (as an illustration for all the
services), including the concrete "HandlerJob" and "Dispatcher" implementations, use a concrete implementation of
"MongoAnnotationClient" to provide the database operations needed by the handler framework. This section provides
further details about how that is accomplished.

The diagram below shows the framework of classes comprising the Data Platform's database interface, again using the
Annotation Service implementation as an example.

MongoClientBase

mongoDatabaseName

init()
fini()

<<Interface>>
MongoAnnotationClientInterface

MongoAsyncClient

mongoCollectionAnnotations

initMongoClient()
initMongoDatabase()
initMongoCollectionAnnotations()
createMongoIndexAnnotations()

MongoSyncClient

mongoCollectionAnnotations

initMongoClient()
initMongoDatabase()
initMongoCollectionAnnotations()
createMongoIndexAnnotations()

MongoSyncAnnotationClient

insertDataSet()
insertAnnotation()
executeQueryAnnotations(request)
findDataSet(dataSetId)

MongoAnnotationHandler

mongoAnnotationClient

newMongoSyncAnnotationHandler()

QueryAnnotationsJob

execute()

AnnotationsResponseDispatcher

handleResult(dbCursor)

The core class in the framework is "MongoClientBase". This abstract class is intended to be extended by intermediate abstract classes that correspond to the various MongoDB Java driver implementations that can be used to provide access to MongoDB operations by the service handler framework. There are at least three different drivers, including "sync", "async", and a new async driver called "reactivestreams".

We wanted to compare the performance of the "sync" and "reactivestreams" drivers in the context of the Data Platform implementations (in particular for the Ingestion Service implementation), so we built the intermediate abstract classes "MongoSyncClient" using the MongoDB "sync" driver and "MongoAsyncClient" using the MongoDB "reactivestreams" driver.

We used this approach, with an abstract generic base class and two intermediate abstract classes, because the way the two MongoDB Java drivers define the same class names (like MongoClient, MongoCollection, MongoDatabase) in different packages ("com.mongodb.client" for the "sync driver and "com.mongodb.reactivestreams.client" for the "reactivestreams" driver).

The base class "MongoClientBase" primarily defines an abstract interface with methods like "initMongoClient()", "initMongoDatabase()", "initMongoCollectionAnnotaitons()", and "createMongoIndexAnnotations()". The two intermediate classes "MongoSyncClient" and "MongoAsyncClient" implement those methods using the "sync" and "reactivestreams" drivers, respectively.

Separately, each service defines an interface whose methods define the database operations needed by the service implementation. For the Annotation Service, the interface "MongoAnnotationClientInterface" defines methods like "insertDataSet()", "insertAnnotation()", "executeQueryAnnotations()" and "findDataSet()". An interface is used so that a concrete interface implementation can derive from either "MongoSyncClient" or "MongoAsyncClient" since the code to insert and retrieve documents is different for the underlying "sync" and "reactivestreams" MongoDB Java drivers.

For the Ingestion Service implementation, we developed both "MongoSyncIngestionClient" and "MongoAsyncIngestionClient" so that we could compare the performance of the two Java drivers in the context of data ingestion. The "sync" driver appears to out-perform the "reactivestreams" driver for our ingestion use case, though we have not performed exhaustive testing and will probably do so at some point. The service implementations are asynchronous by design, so we don't need to use an asynchronous MongoDB Java driver to accomplish that objective.

For the Query and Annotation Services, where performance is important but not on the same level as the Ingestion Service, we built a single concrete database client implementation extending the intermediate class "MongoSyncClient". This class inherits the database client connection management framework from "MongoClientBase", and provides overrides for the service-specific interface methods in "MongoAnnotationClientInterface" to provide the database operations needed by the service implementation.

The handler for each service implementation including "MongoAnnotationHandler" provides factory methods such as "newMongoSyncAnnotationHandler()" to create a handler instance initialized with the synchronous database client "MongoSyncIngestionClient". When new jobs are added to the handler's task queue, they are provided with a reference to the database client for accessing database operations by both the job and its dispatcher, as needed. This is illustrated in the diagram by "QueryAnnotationsJob" and the corresponding "AnnotationsResponseDispatcher" which call the database client methods "executeQueryAnnotations()" and "findDataSet()" in the execution of their methods "execute()" and "handleResult()", respectively.

serialization of protobuf objects to MongoDB documents

The Ingestion Service adds a document to the MongoDB "buckets" collection for each "DataColumn" (vector of samples) contained in an ingestion request's data frame. The documents contain the vector's "DataValues" and the corresponding "DataTimestamps" for those values, as well as other details such as bucket start and end time.

Originally, the Ingestion Service used a polymorphic hierarchy of parameterized classes derived from the base class, "BucketDocument" to handle the various data types supported by the API "DataValue" message. This approach required a Java "POJO" class for each supported API "DataValue" type for mapping the API type to a Java type. For example, the Java class "DoubleBucketDocument" extends the base "BucketDocument" class with the type parameter "Double" to handle the API DataValue value type "doubleValue", and the Java class "StringBucketDocument" with parameter type "String" handles the API DataValue value type "stringValue".

While this approach is fairly straightforward for the basic scalar data types defined by the API that map to primitive Java types, it encounters problems handling more complex API data types like "Array" and "Structure" which can contain nested references to other "DataValue" types. It would be challenging to design Java BucketDocument subclasses to handle arbitrarily nested arrays of structures that might contain arrays of different data types, etc.

In order to support these more complex API data types, we changed approach in version 1.4 to serialize the protobuf "DataColumn" object directly to the "BucketDocument" as an opaque byte array field without otherwise unpacking the column's "DataValues". This allows us to support all API data types (including arbitrarily nested arrays and structures) while eliminating the need for the hierarchy of classes derived from "BucketDocument" for handling each unique API data type.

In addition to simplifying the code base, this change improved ingestion performance significantly and reduced the MongoDB storage footprint (due to the compression used in protobuf serialization).

We decided to use the same approach for storing "DataTimestamps", which can contain either a "SamplingClock" (specifying the start time, number of samples, and sample period) or "TimestampsList" (with an explicit list of data timestamps), in the "BucketDocuments". Instead of unpacking the "DataTimestamps" from the ingestion request and adding conditional fields to "BucketDocument" to store the constituent "SamplingClock" or "TimestampsList", we simply serialize the "DataTimestamps" protobuf object to the "BucketDocument" as an opaque byte array field.

The class diagram below shows part of the original polymorphic "BucketDocument" hierarchy. The base class uses the type parameter "T" to indicate the type of data managed by the class. It also includes the "@BsonDiscriminator" annotation to specify that MongoDB should use the value of the "dataType" field to determine which Java POJO class to use for each document. The derived class "DoubleBucketDocument" uses the type parameter "Double" to show that the bucket contains Java Double data values. Similarly, the "StringBucketDocument" derived class manages "String" data values.

The diagram also shows that the base class uses fields to capture the details for the bucket's "DataTimestamps", with fields for both "SamplingClock" and "TimestampsList" details.
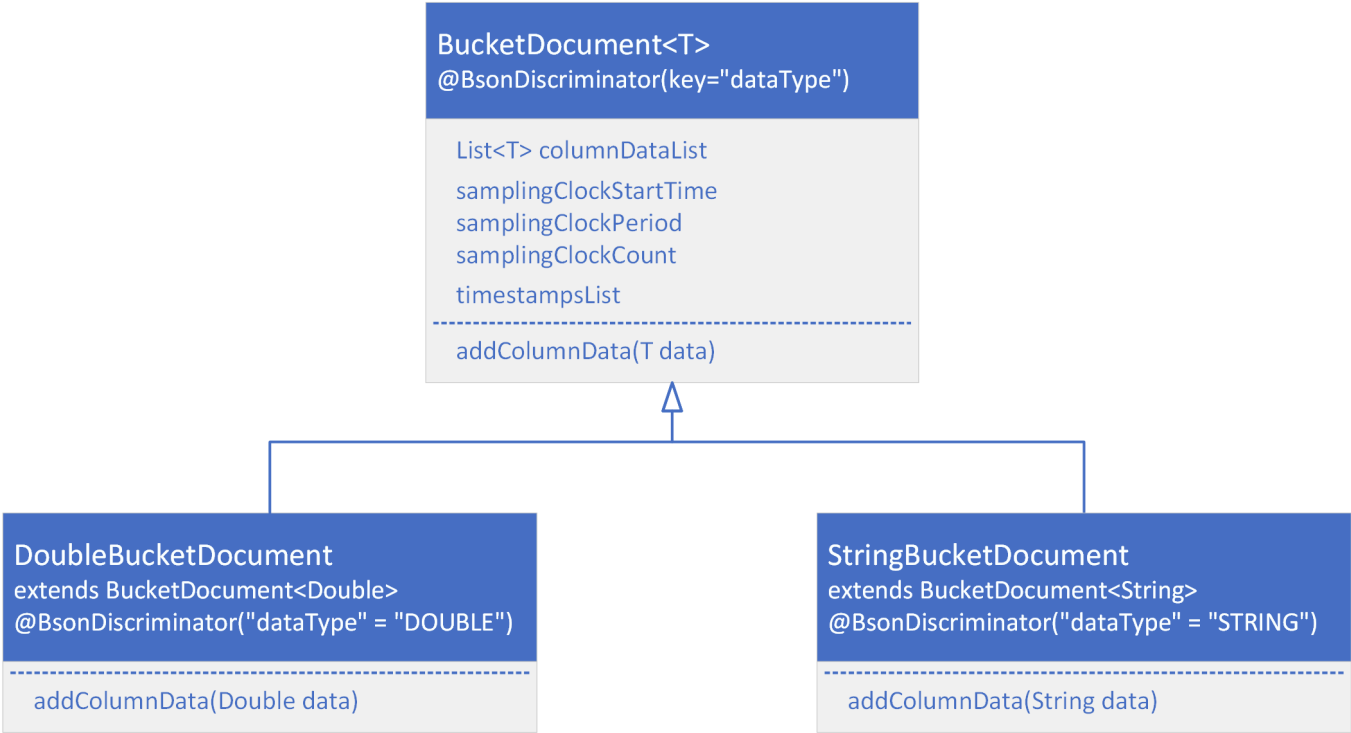
In Data Platform version 1.4, the polymorphic "BucketDocument" hierarchy shown above is replaced by the new standalone "BucketDocument" class shown in the class diagram below. It contains byte array fields "dataColumnBytes" and "dataTimestampsBytes" to hold the serialized "DataColumn" and "DataTimestamps" for the bucket, respectively. It provides methods for reading and writing the serialized object to the bucket, and getting the API type case and name for the serialized object.

BucketDocument

byte[] dataColumnBytes
int dataTypeCase
String dataType
byte[] dataTimestampsBytes
int dataTimestampsCase
String dataTimestampsType

----

writeDataColumnContent(DataColumn)
readDataColumnContent()
getDataTypeCase()
getDataType()
writeDataTimestampsContent(DataTimestamps)
readDataTimestampsContent()
getDataTimestampsCase()
getDataTimestampsType()

configuration

The Data Platform service implementations all share a simple configuration framework. The primary objective for the configuration framework is to minimize the code required to retrieve configuration values, such as checking if the resource is defined, checking for a null return value, providing a default value, and casting to common Java types. We wanted to define methods on the configuration tool that hide those details from the caller as much as possible.

The primary configuration mechanism is a "yaml" file named "application.yml". An override file name can be specified either on the command line (using "dp.config") or as an environment variable (using "DP.CONFIG"). Individual config resources can be overridden on the command line by prefixing them with "dp.".

The keys in the configuration file are hierarchical using "." notation. Given the config file content below:
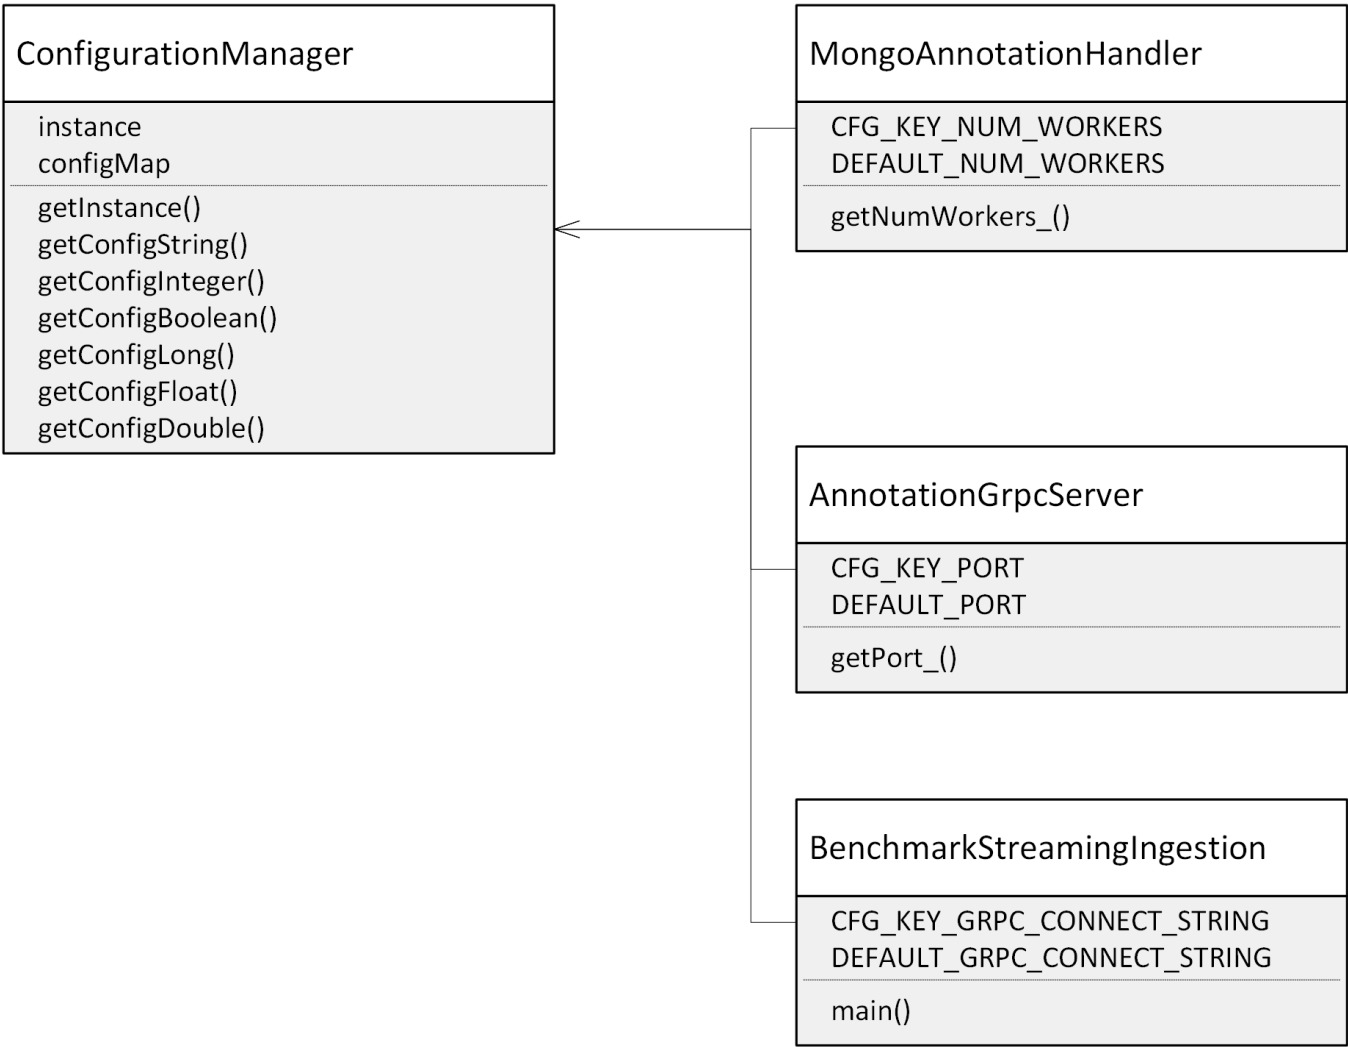
 # annotationHandler: Settings for the annotation Service request handler.

annotationHandler:

_ # annotationHandler.numWorkers: Number of worker threads used by the annotation Service request handler._

_ # This parameter might take some tuning on deployments to get the best performance._

_ numWorkers: 7_

the number of workers for the annotation service is obtained using the key "AnnotationHandler.numWorkers".

The diagram below shows the Data Platform "ConfgurationManager", with some example uses including "MongoAnnotationHandler", "AnnotationGrpcServer", and "BenchmarkStreamingIngestion", which retrieve number of workers, port number, and connect string, respectively.

ConfigurationManager
instance configMap
getInstance() getConfigString() getConfigInteger() getConfigBoolean() getConfigLong() getConfigFloat() getConfigDouble()

MongoAnnotationHandler
CFG_KEY_NUM_WORKERS DEFAULT_NUM_WORKERS
getNumWorkers_()

AnnotationGrpcServer
CFG_KEY_PORT DEFAULT_PORT
getPort_()

BenchmarkStreamingIngestion
CFG_KEY_GRPC_CONNECT_STRING DEFAULT_GRPC_CONNECT_STRING
main()

The "ConfigurationManager" provides methods for accessing config resources casted to common Java types such as "getConfigString()" and "getConfigInteger()". Each accessor method has two variants, one that returns the casted resource value of null if not defined, and the other that uses a default value parameter to return the default value if the resource is not defined. This allows the caller to not check for null, not cast the value, and not override with a default, etc.

The "ConfigurationManager" is a "singleton" pattern implementation, so the first call to its "getInstance()" method initializes the config manager by reading the default or specified override config file and applying any overrides of

individual resource values. Most usage within Data Platform applications looks like the snippet below:

configMgr().getConfigInteger(CFG_KEY_NUM_WORKERS, DEFAULT_NUM_WORKERS)

where "configMgr()" is a convenience method for accessing the singleton "ConfigurationManager" instance.

exporting data

Export API requests are handled in the same fashion as all other API requests to the Data Platform Services. A job is created for handling the request and added to the handler's task queue. When the job is executed, it 1) retrieves the associated dataset from the database, 2) generates output file paths and URLs, 3) creates the directories in the file path as neeeded, and 4) exports the data for the dataset to file.

The diagram below shows the classes involved in the export handling framework.
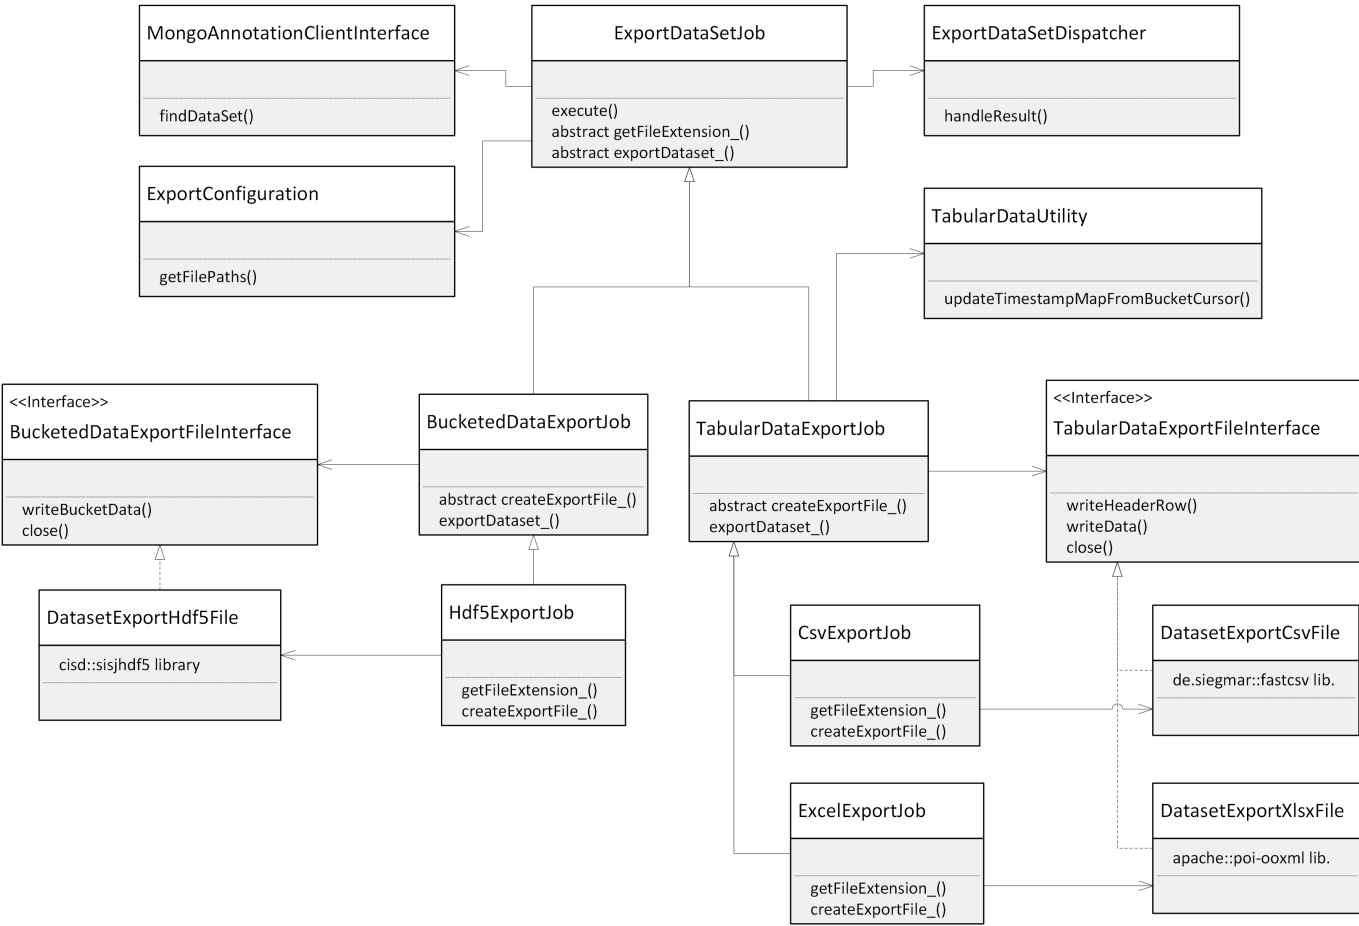
There is a single base class for all export jobs, ExportDataSetJob that controls the execution. It defines abstract methods that are overridden by derived classes to handle writing data to the particular file format. It uses the MongoAnnotationClientInterface to find the specified dataset. It uses ExportDataSetDispatcher to send a response in the API method's response stream.

The new class ExportConfiguration is responsible for reading the export-related config resources and generating file paths and URLs based on that configuration, and is used by the job to generate paths for the export file.

There are two intermediate classes, BucketedDataExportJob and TabularDataExportJob, that provide common logic for writing time-series data organized in "buckets" and writing tabular file formats, respectively.

There is a single concrete class for writing bucketed data, Hdf5ExportJob, that handles writing bucketed data to HDF5 format.

TabularDataExportJob uses utility methods in TabularDataUtility to build a data structure for iterating and accessing data in a tabular fashion. This functionality is contained in a standalone class so that it can be shared with other code working with tabular data (such as building the result for a tabular data query).

There are two concrete classes for writing tabular data, CsvExportJob and ExcelExportJob, that handle writing tabular data to CSV and XLSX formats, respectively.

The details of writing to specific file formats from the bucketed and tabular data export job classes are defined by two interfaces, BucketedDataExportFileInterface and TabularDataExportFileInterface, respectively. Hdf5ExportJob uses the concrete interface implementation, DatasetExportHdf5File to write data to HDF5 format. CsvExportJob uses DatasetExportCsvFile to write to CSV files. ExcelExportJob uses DatasetExportXlsxFile to write to XLSX files.

The file interface implementation classes use new third party library dependencies (defined in "pom.xml" for lower-level file access APIs. The libraries "sis-base" and "sis-jhdf5" from "cisd" are used by DatasetExportHdf5File for writing HDF5 files. The library "fastcsv" from "de.siegmar" is used by DatasetExportCsvFile for writing CSV files. The library "poi-ooxml" from "org.apache.poi" is used by DatasetExportXlsxFile for writing XLSX files.

performance benchmarking

Because performance is the most important requirement for the Data Platform Ingestion Service, we developed a benchmarking framework in parallel with the service implementation so that we could measure performance at each stage of development and compare different approaches (such as comparing the performance obtained with the MongoDB Java "sync" driver with the "reactivestreams" one). The same pattern was followed to build a performance benchmark framework for the Query Service implementation.

Both benchmark frameworks use the MongoDB database "dp-benchmark". Before each run of any benchmark, the contents of that database are removed and new contents are added by the benchmark.

Benchmark-specific servers, "BenchmarkIngestionGrpcServer" and "BenchmarkQueryGrpcServer", were added that override the standard Ingestion and Query service network ports for those services to avoid adding benchmark data to a live production database.

ingestion service performance benchmarking

The diagram below shows the elements of the ingestion performance benchmarking framework.
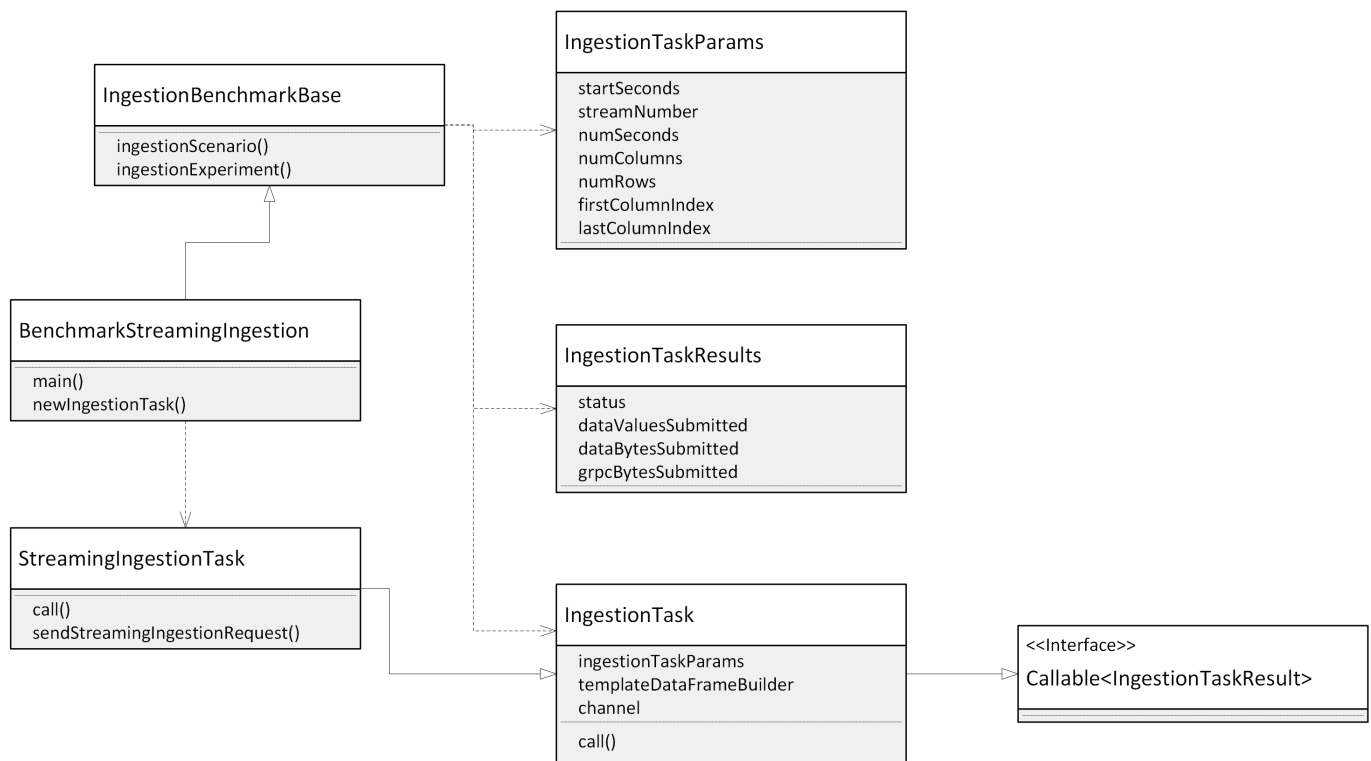
The framework supports running ingestion scenarios for different approaches, so the base class "IngestionBenchmarkBase" contains the key framework components for running an ingestion benchmark. It defines the nested class "IngestionTask" to encapsulate the logic for invoking an ingestion API, creating the stream of ingestion requests, and handling the API response stream. The nested classes "IngestionTaskParams" and "IngestionTaskResults" are used to contain the parameters needed by the task and to return performance results from the task.

"IngestionBenchmarkBase" provides the method "ingestionExperiment()" for sweeping combinations of parameter values for variables like number of threads and number of API streams, and calculating an overall performance benchmark for the run. The lower level method "ingestionScenario()" is used by the experiment driver method to run an individual scenario and measure its performance. This method is also used in the integration test framework to create a regression test that not only runs the ingestion scenario to create data in the archive, but verifies database contents and API responses. See section "integration testing" for more details.

The application class "BenchmarkStreamingIngestion" extends the base class to run a performance benchmark for the "ingestDataStream()" bidirectional streaming Ingestion Service API. It defines the nested class "StreamingIngestionTask", implementing the "call()" method to call the API, send a stream of requests, handle the response stream, and collect performance stats.

query service performance benchmarking

The query service benchmark framework is a bit more complicated than the ingestion service framework because it also loads the data into MongoDB to be used in the performance benchmark. Initially, it utilized data loaded by the ingestion benchmark, but we decided to make it a standalone application and added the data-loading mechanism. The diagram below shows the query service benchmark framework.
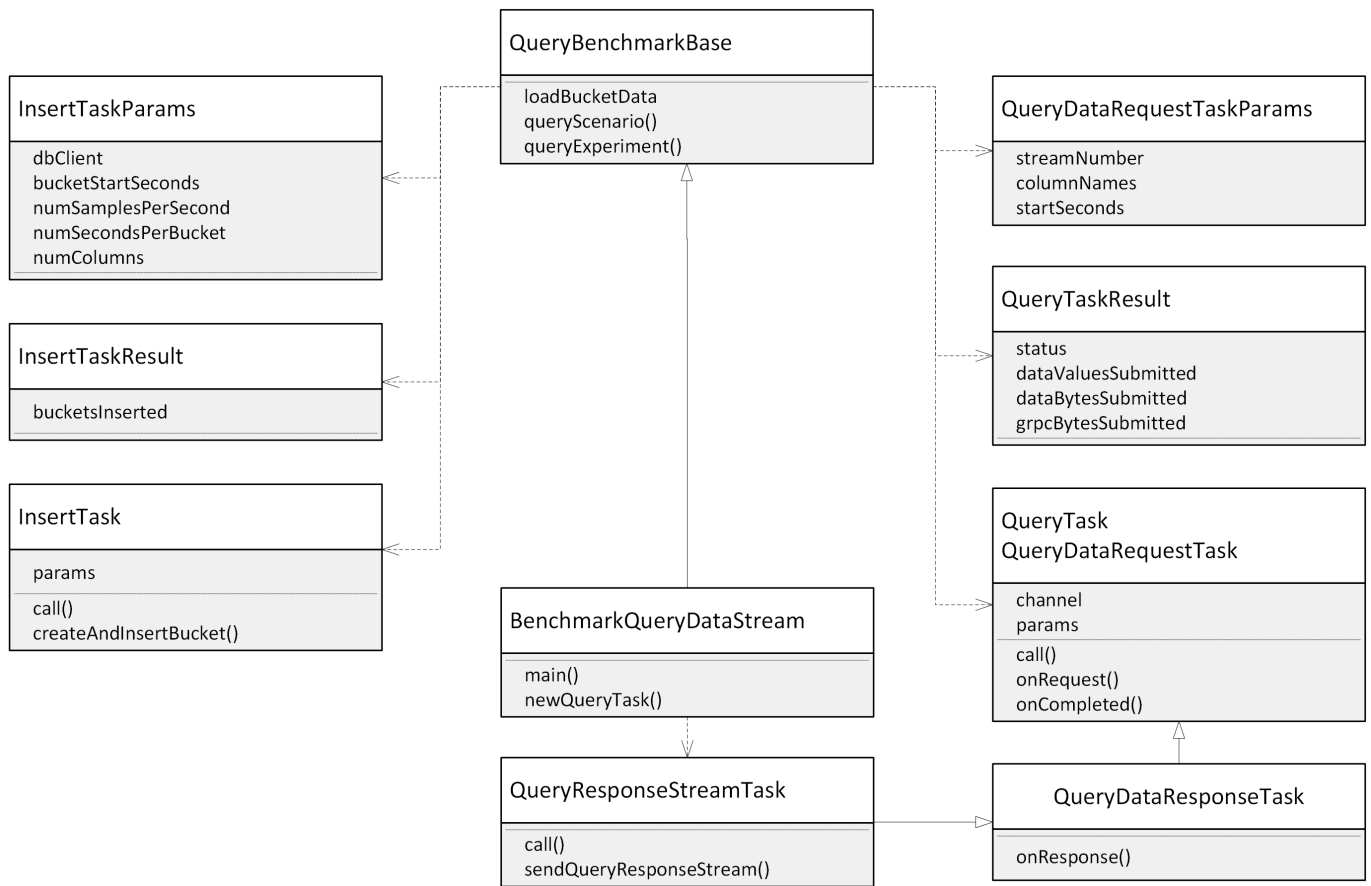
The core class of the framework is "QueryBenchmarkBase". For loading data, it uses a multithreaded "ExecutorService" with the custom "Callable" task class "InsertTask". Each task loads data as specified by the "InsertTaskParams" passed to the task, and returns a "InsertTaskResult" with details including the number of data buckets inserted by the task. Loading data is triggered by the method "loadBucketData()".

The query performance benchmark part of the framework follows a similar pattern as described for the ingestion benchmark, above. The base class provides a higher-level method "queryExperiment()" to sweep parameter value combinations for total number of PVs, number of PVs per request, and number of threads. The experiment driver uses the lower-level method "queryScenario()" to measure performance for a particular combination of parameter values.

An "ExecutorService" is used with custom tasks that extend the classes "QueryTask", "QueryDataRequestTask", or "QueryDataResponseTask" to call a specific query API, handle the response stream, and calculate statistics.

Various concrete performance benchmark application classes extend the base class, each measuring the performance of a particular query API. The query benchmark application classes include "BenchmarkQueryDataStream", "BenchmarkQueryDataBidiStream", and "BenchmarkQueryDataUnary". Each of them defines a concrete custom task class that extends one of the framework task base classes.

For example, the benchmark application class "BenchmarkQueryDataStream" shown in the diagram above defines the task class "QueryResponseStreamTask" that extends "QueryDataResponseTask". Each task calls the "queryDataStream()" API and keeps track of the number of values and bytes sent for use in performance statistics.

generating sample data

Recognizing the previous use of the ingestion performance benchmark application for creating test data for web app development and demo, a new TestDataGenerator utility has been added. A more fully featured simulator / generator is under development, but in the meantime this tool can be used to generate sample data for use in web application development or demo purposes. Like the ingestion benchmark, it generates one minute's data for 4000 PVs sampled at 1 KHz.

To use the sample data generator, the standard ingestion service should be running ("IngestionGrpcServer"). This captures data to the standard "dp" database instead of the benchmark-specific "dp-benchmark" database.

The dp-support repo contains a wrapper script in the bin directory for running the sample data generator, "app-run-test-data-generator".

regression testing

The primary objective for the regression test suite was to allow coverage to be added for essentially any part of the Data Platform common code and service implementations including lower level components, and for the most part that has been accomplished. Test coverage for the Ingestion and Query service implementations is pretty extensive, and covers some of the lower-level features that are hard to cover in higher-level scenarios.

All regression tests using a MongoDB database named "dp-test". The Data Platform's MongoDB schema is described in Section "dp-service MongoDB schema and data flow".

After adding coverage for both the Ingestion and Query service implementations, we added an "integration testing" framework that provides a mechanism for running higher-level scenarios that involve any/all of the service implementations. That framework is discussed in section "integration testing".

Since the integration testing framework was added, we've preferred adding test coverage at that level when possible because it exercises the communication framework in addition to the service implementations. For that reason, there is less low-level test coverage of the annotation service implementation than the other service, but pretty good coverage at the higher-level integration test level. We will add more extensive low-level coverage for all the services as time goes on to cover more special / unusual cases.

We've used a naming convention for classes that contain jUnit test cases to end the class name with "Test". The names of base classes that are intended to be extended by concrete test classes end with "Base". The base and test classes are summarized below.

ConfigurationManager tests (com.ospreydcs.dp.service.common.config)

class	description
ConfigurationManagerTestBase	Base class for ConfigurationManager test coverage. Provides nested class ConfigurationManagerDerived that allows singleton instance to be reset during test execution. Provides utility methods for working with an override config file.
ConfigurationManagerTest	Includes test01GetConfig() for testing the ConfigurationManager getter methods without default values, and test02GetConfigWithDefault() for testing

	the getter methods with default values. test03GetInstance() tests the singleton access method.
ConfigurationManagerOverrideConfigFileCmdLineTest	Provides coverage for overriding config file via command line.
ConfigurationManagerOverrideConfigFileEnvVarTest	Provides coverage for overriding config file via environment variable.
ConfigurationManagerOverridePropertyCmdLineTest	Provides coverage for overriding individual config resources via command line.

mongo tests (com.ospreydcs.dp.service.common.mongo)

class	description
MongoTestClient	Extends "MongoSyncClient" database framework class to provide database-level utilities to regression tests. The "init()" method drops the test database if it exists and creates a new one. Provides utilities useful for verifying database artifacts including "findBucket()", "findRequestStatus()", "findDataSet()", and "findAnnotation()".
MongoSyncDriverTest	Used to confirm expected behavior of MongoDB Java "sync" driver, independent of the Data Platform database client framework for special cases such as the behavior when a duplicate database id is inserted, etc.
MongoAsyncDriverTest	Used to confirm expected behavior of MongoDB Java "reactivestreams" driver, independent of the Data Platform database client framework for special cases such as the behavior when a duplicate database id is inserted, etc.
MongoIndexTest	This simple test checks that the expected indexes exist on each of the Data Platform MongoDB collections. The indexes are created at startup by MongoClientBase. For each collection, there is a section in the test that checks for the expected number of indexes, and checks each of the expected index names (using the standard Mongo index naming convention).

ingestion service tests (com.ospreydcs.dp.service.ingest)

class	description
IngestionTestBase	Base class for derived test classes providing ingestion service test coverage. Includes nested utility classes "IngestionRequestParams" to encapsulate the parameters for an ingestion request and "IngestionResponseObserver" for handling the response stream from the streaming ingestion API. Includes utility method variants of "buildIngestionRequest()" for building an ingestion request from a params object.
MongoIngestionHandlerTestBase	Provides coverage for ingestion service handler framework, with detailed positive and negative test cases. This is a base class so that it can be

	extended by two concrete classes, one each for testing the sync and async database client framework classes.
MongoSyncIngestionHandlerTest	Extends MongoIngestionHandlerTestBase to run positive and negative ingestion service test cases using the MongoSyncIngestionClient database client for database operations.
MongoAsyncIngestionHandlerTest	Extends MongoIngestionHandlerTestBase to run positive and negative ingestion service test cases using the MongoAsyncIngestionClient database client for database operations.
IngestionValidationUtilityTest	Provides coverage for various rejection scenarios in IngestionValidationUtility.validateIngestionRequest().
IngestionServiceImplTest	Provides coverage for utility methods in IngestionServiceImpl.

query service tests (com.ospreydcs.dp.service.query)

class	description
QueryTestBase	Base class for derived test classes providing query service test coverage. Includes nested classes "QueryDataRequestParams" to encapsulate parameters for a time-series data query request, "QueryTableRequestParams" to encapsulate parameters for a table-formatted time-series data query request, and "QueryResponseTableObserver" / "QueryResponseStreamObserver" / "QueryMetadataResponseObserver" for handling the response stream from the table query, time-series data query, and metadata query APIs, respectively. Provides utility methods for building gRPC request objects from params objects.
QueryGrpcTest	Provides basic coverage of query service gRPC communication using the in-process gRPC framework.
QueryServiceImplTest	Provides test coverage for QueryServiceImpl utility methods that generate gRPC responses.
QueryHandlerUtilityTest	Provides coverage for validation utility methods in QueryHandlerUtility.
MongoQueryHandlerTestBase	Provides coverage for query service handler framework, with detailed positive and negative test cases. This is a base class so that it can be extended by two concrete classes, one each for testing the sync and async database client framework classes.
MongoSyncQueryHandlerTest	Extends MongoQueryHandlerTestBase to run positive and negative ingestion service test cases using the MongoSyncQueryClient database client for database operations.
MongoQueryHandlerErrorTest	Provides coverage for an error response to query request by returning a null cursor in response to a database query, which simulates an error at the MongoDB database level (as opposed to the Query Service application level).

annotation service tests (com.ospreydcs.dp.service.annotation)

class	description
AnnotationTestBase	Base class for derived test classes providing annotation service test coverage. Includes nested classes "CreateDataSetParams", "CreateAnnotationRequestParams", "CreateCommentAnnotationParams", "AnnotationDataSet" and "AnnotationDataBlock" to encapsulate the parameters for Annotation Service API requests. Provides utility methods "buildCreateCommentAnnotationRequest()" and "buildQueryAnnotationsRequestOwnerComment()" for building gRPC request objects from params objects for the create annotation and query annotations APIs, respectively.
ExportConfigurationTest	Provides coverage for generating file paths and URL from export-related config resources. Includes positive test cases for full and minimal configuration handling, and negative test cases for invalid configurations.
DatasetExportHdf5FileTest	Provides unit test coverage for writing a DataSetDocument object to HDF5 file format using class DatasetExportHdf5File.

integration testing

One requirement for this project is to provide integration testing that provides coverage for scenarios that involve multiple services. We developed an integration testing framework that supports creating tests that include data ingestion, query, and annotation. The framework is shown in the diagram below.

The core of the framework is the class "GrpcIntegrationTestBase". The base class is extended by test classes such as "StaggeredTimestampTest", "MetadataQueryTest", "TableQueryTest", and "AnnotationTest". There are several integration tests focused on different aspects of the Ingestion Service that are not shown in the diagram, but they follow the same pattern (inheriting from "GrpcIntegrationTestBase" and using utility methods in the base class and "IngestionTestBase"). These are discussed in the table summary of the integration tests below. A special case is "BenchmarkIntegrationTest", which is detailed in the next subsection.

The framework is built using the "in-process" gRPC framework, which allows multiple gRPC clients and servers to run in a single Java process. It is a convenient way to build integration test coverage without dealing with synchronizing multiple processes.

The base class's "setUp()" method brings up the Data Platform environment, including Ingestion, Query, and Annotation services using the corresponding service implementation's "init()" method (e.g., "IngestionServiceImpl.init()"). It creates a gRPC channel for each service for invoking RPC methods on that service. The base class "tearDown()" method shuts down the services using each implementation's "fini()" method.

The base class provides convenience methods for calling service API's, validating the corresponding database artifacts, and verifying the API response stream. For example, the "AnnotationTest" uses the following base methods to 1) create data in the archive, 2) create data blocks, data sets, and annotations, 3) query annotations, and 4) query time-series data using the data blocks returned by the annotations query:

1. ingestDataStreamFromColumn()

2. sendAndVerifyCreateDataSet()
3. sendAndVerifyCreateCommentAnnotation()
4. sendAndVerifyQueryAnnotationsOwnerComment()
5. queryDataStream()

Each service implementation provides a test base class with data structures, classes, and utilities for using the service APIs. For example, "IngestionTestBase" includes the data structure "IngestionRequestParams" to contain the parameters for an ingestion request. The method "buildIngestionRequest()" creates a gRPC ingestion request object from the params object. The nested class "IngestionResponseObserver" is a response stream observer for the streaming ingestion API that builds a list of replies for use in verifying the API response. The classes "QueryTestBase" and "AnnotationTestBase" offer similar utilities for the Query and Annotation Services, respectively.

For API calls that write data to the database, the base class wrapper methods validate the corresponding database artifacts for each API request. The class "MongoTestClient" provides utilities for retrieving items from the database for validation. For example, the base class method "ingestDataStreamFromColumn()" calls the bidirectional streaming ingestion API to create data, and then uses "MongoTestClient.findBucket()" and "findRequestStatus()" to retrieve the corresponding database artifacts, comparing their contents to the results expected for the request. Similarly, the base class methods "sendAndVerifyCreateDataSet()" and "sendAndVerifyCreate…Annotation…()" using test client methods "findDataSet()" and "findAnnotation()", respectively.

The various integration test classes are summarized in the tables below.

com.ospreydcs.dp.service.integration.annotation

class	description
AnnotationTest	This test provides coverage for various aspects of the AnnotationService API including creating and querying DataSets, and creating and querying Annotations. The test runs a simple ingestion scenario to create data for various devices. It includes both negative and positive tests for createDataSet() using the ingested data. It includes both negative and postive tests for the queryDataSets() API over the DataSets created by the test. It includes both negative and positive tests for createAnnotation() using the DataSets created by the test. It includes negative and positive coverage for queryAnnotations() using the annotations created by the test. It includes a scenario to run a time-series data query using the DataSet returned by queryAnnotations(). It includes positive and negative test cases for exporing datasets to HDF5, CSV, and XLSX formats.

com.ospreydcs.dp.service.integration.ingest

class	description
DataTypesTestBase	This base class includes test scenarios for ingestion of "complex" data types, including 2D array, image, and structure data and verification of ingestion results. It defines the abstract method sendAndVerifyIngestionRpc_() that is overridden by subclasses to call one of the ingestion APIs.
DataTypesUnaryTest	Extends DataTypesTestBase and overrides sendAndVerifyIngestionRpc_() to test ingestion of complex data types using the unary ingestion API, ingestData().

DataTypesStreamingTest	Extends DataTypesTestBase and overrides sendAndVerifyIngestionRpc_() to test ingestion of complex data types using the bidirectional streaming ingestion API, ingestDataStream().
ExplicitTimestampListTest	Provides coverage for ingestion of data using DataTimestamps with an explicit TimestampsList (instead of SamplingClock, covered in most of the other ingestion integration tests).
ReisterProviderTest	Provides coverage for the registerProvider() API method.
RequestStatusTest	Provides coverage for the queryRequestStatus() API method.
UnaryTest	Provides coverage for ingestion of data using the unary ingestion API, ingestData(). Includes simple negative and positive test scenarios.
UnidirectionalStreamTest	Provides coverage for the ingestDataStream() API method.
ValidationTest	Currently sort of a placeholder, but provides coverage for a single rejection scenario. There is more extensive rejection test coverage in the regular unit test IngestionValidationUtilityTest, but I wanted a rejection test that exercises the gRPC communication framework to make sure it works as expected at the API level.

**com.ospreydcs.dp.service.integration.query**

class	description
MetadataQueryTest	This test provides coverage for the Query Service's queryMetadata() API. It runs a simple ingestion scenario, and then various negative and positive test cases for queryMetadata().
StaggeredTimestampTest	This test provides coverage for queries in a more interesting scenario for ingested data. I added it when we added the queryDataTable() API to make sure that the query table result was correct when ingested PV data used different sample periods, and when the query time range was intentionally offset from the data bucket begin times.
TableQueryTest	This test provides coverage for the queryDataTable() API using a simpler ingestion scenario than StaggeredTimestampTest. It includes positive test cases for both column and row-oriented query table result.

benchmark integration test

The integration test class "BenchmarkIntegrationTest" is a special case, combining the integration testing and performance benchmarking frameworks in a single test case. This test runs an ingestion scenario for a larger universe of data, including one minute's data for 4,000 PVs each sample 1,000 times per second. It then exercises various time-series data queries against that data. The test validates the database artifacts created by the ingestion scenario, and verifies all API responses including both ingestion and query.

By exercising the streaming data ingestion API method and the three time-series data query API methods, this test covers a large part of the codebase for the Data Platform Ingestion and Query Services.

The benchmark integration test framework is shown in the diagram below.
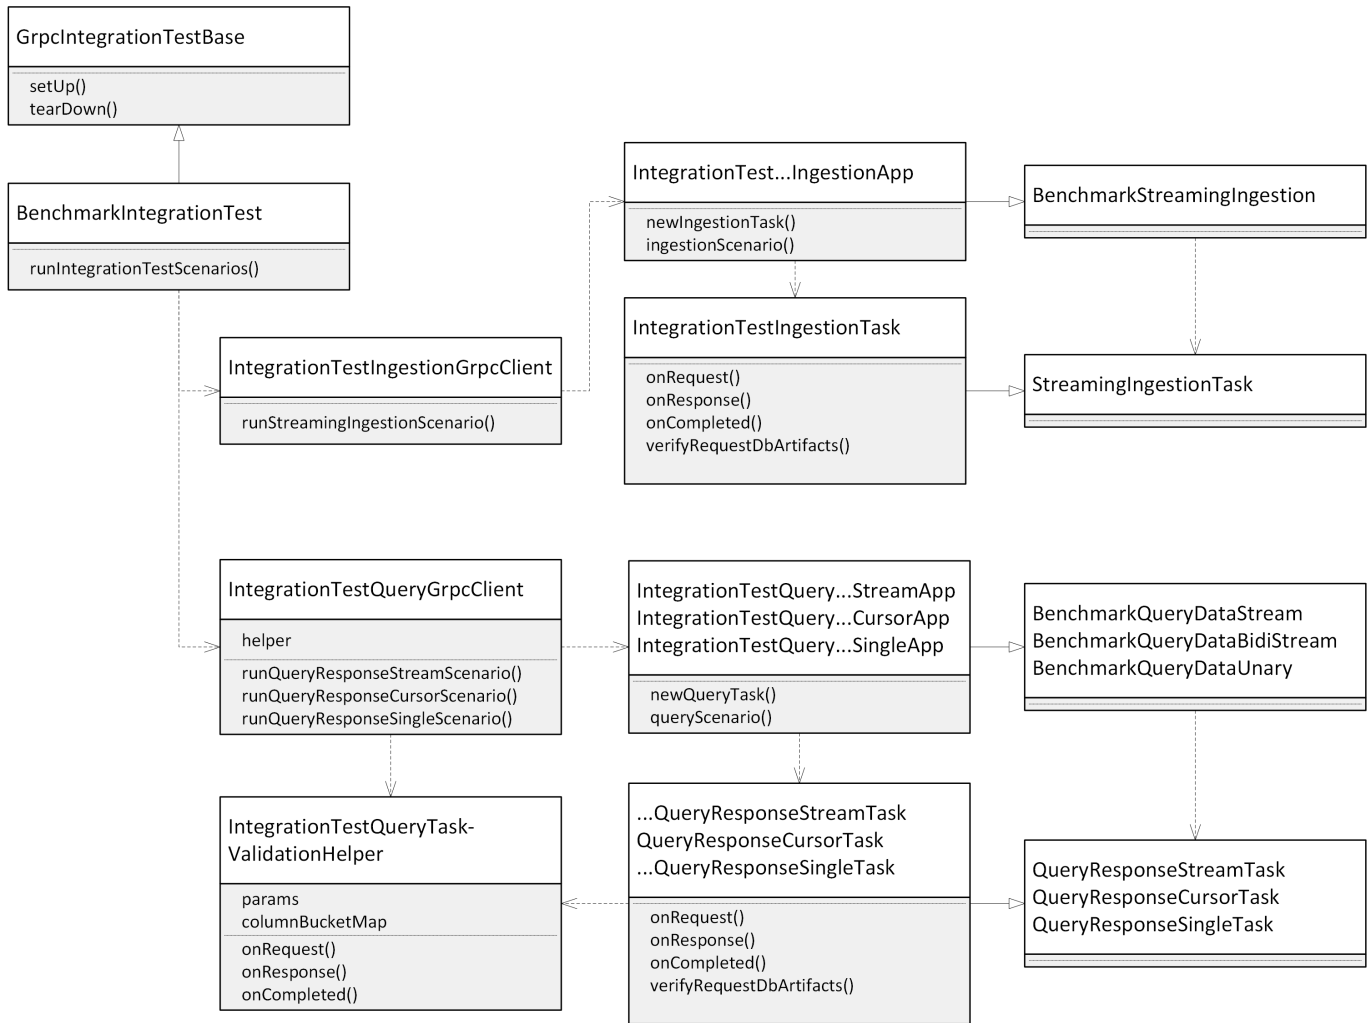
Like the other integration tests, the benchmark integration test extends "GrpcIntegrationTestBase" and uses "setUp()" and "tearDown()" to bring up the Data Platform services in a single Java process using the in-process gRPC framework.

Unlike the other tests, however, it does not use the base class's wrapper methods for invoking API methods and verifying their results. Instead, it uses the benchmark framework for that purpose.

The "BenchmarkIntegrationTest" defines the class "IntegrationTestStreamingIngestionApp" which extends the benchmark framework class "BenchmarkStreamingIngestion". The ingestion app class defines an extension to "StreamingIngestionTask" called "IntegrationTestIngestionTask". The latter overrides the validation hook methods "onRequest()", "onResponse()", and "onCompleted()" to build a data structure for validating ingestion requests, and validate the requests on their completion (including both database artifacts and API responses).

The benchmark integration test class uses the nested class "IntegrationTestIngestionGrpcClient" to run an ingestion scenario using the method "runStreamingIngestionScenario()" via the benchmark framework method "IntegrationTestStreamingIngestionApp.ingestionScenario()". This runs the ingestion scenario and triggers the validation hook methods as the scenario is executed.

The same pattern is used to run the query scenarios, though there are more classes involved because we are calling three different query API methods, "queryDataStream()", "queryDataBidiStream()", and "queryData()".

The "queryDataStream()" scenario is handled by the benchmark integration test's nested class "IntegrationTestQueryResponseStreamApp" and its concrete task extension "IntegrationTestQueryResponseStreamTask".

The "queryDataBidiStream()" scenario is handled by the nested class "IntegrationTestQueryResponseCursorApp" with task subclass "IntegrationTestQueryResponseCursorTask".
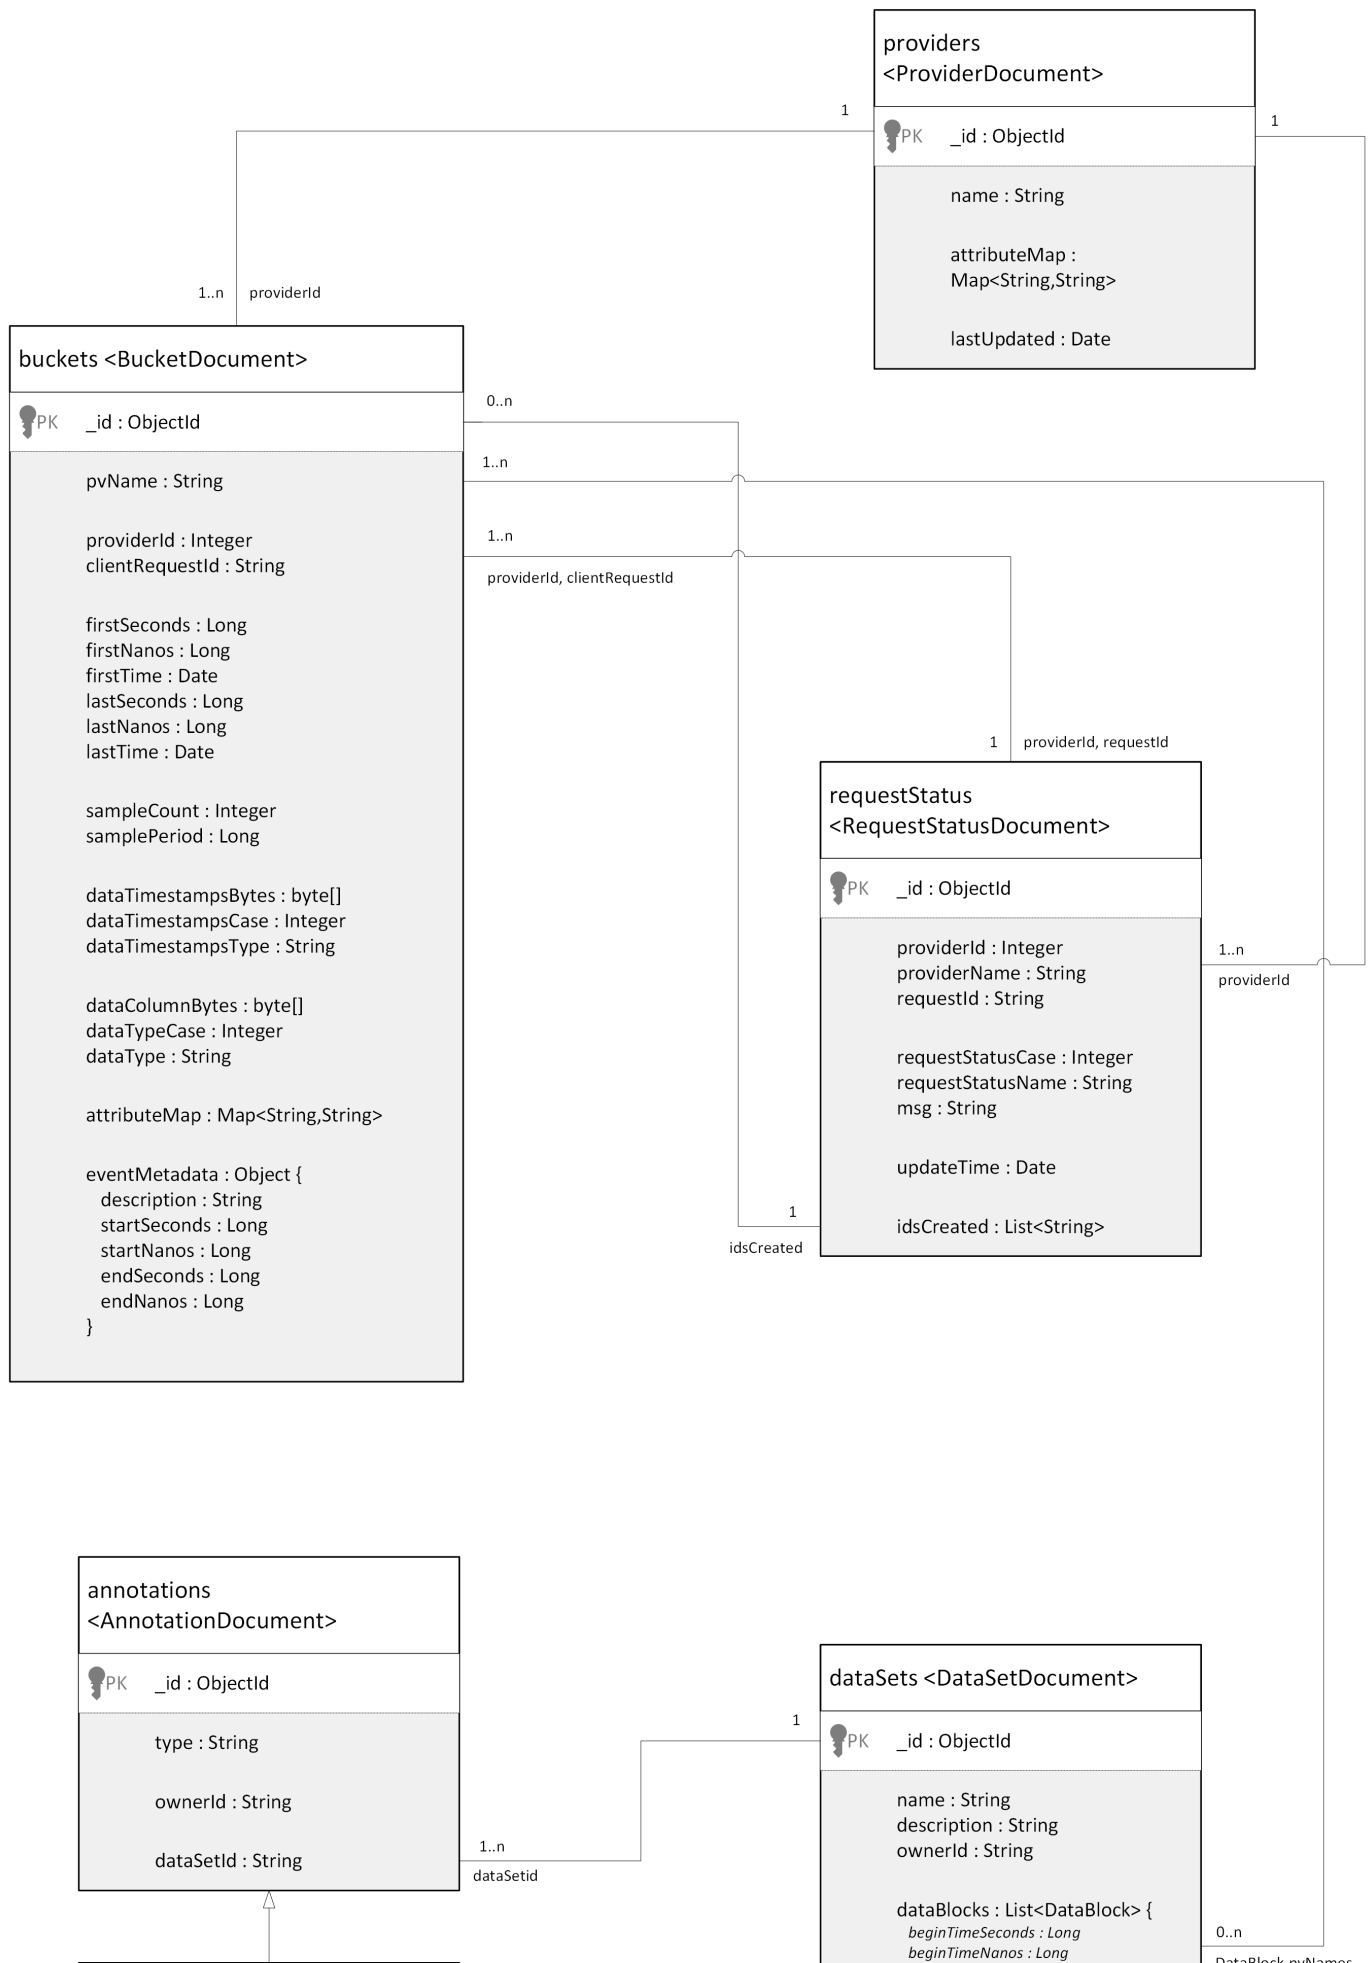
The "queryData()" scenario is handled by nested class "IntegrationTestQueryResponseSingleApp" with task subclass "IntegrationTestQueryResponseSingleTask".

In all three cases, the task subclass dispatches the validation hook methods "onRequest()", "onResponse()", and "onCompleted()" to the "IntegrationTestQueryTaskValidationHelper" managed by the nested class "IntegrationTestQueryGrpcClient". The helper implements the validation hooks to build data structures with information about the API requests, and uses them to verify the query results returned in the response stream.

The "IntegrationTestQueryGrpcClient" provides methods for running each of the three query API scenarios, "runQueryResponseStreamScenario()", "runQueryResponseCursorScenario()", and "runQueryResponseSingleScenario()". Each query scenario retrieves one minute's data for 1,000 PVs, and uses the validation hook method implementations to verify the query results.

dp-service MongoDB schema and data flow

The Data Platform services utilize MongoDB for persistence. The default database is called "dp". Regression tests use a database called "dp-test" whose contents are removed at the start of each test. The diagram below shows the entity-relationship model for the Data Platform MongoDB schema.

providers
<ProviderDocument>

🔑PK   _id : ObjectId

name : String

attributeMap :
Map<String,String>

lastUpdated : Date

---

buckets <BucketDocument>

🔑PK   _id : ObjectId

pvName : String

providerId : Integer
clientRequestId : String

firstSeconds : Long
firstNanos : Long
firstTime : Date
lastSeconds : Long
lastNanos : Long
lastTime : Date

sampleCount : Integer
samplePeriod : Long

dataTimestampsBytes : byte[]
dataTimestampsCase : Integer
dataTimestampsType : String

dataColumnBytes : byte[]
dataTypeCase : Integer
dataType : String

attributeMap : Map<String,String>

eventMetadata : Object {
  description : String
  startSeconds : Long
  startNanos : Long
  endSeconds : Long
  endNanos : Long
}

1..n   providerId

0..n

1..n

1..n

providerId, clientRequestId

---

requestStatus
<RequestStatusDocument>

🔑PK   _id : ObjectId

providerId : Integer
providerName : String
requestId : String

requestStatusCase : Integer
requestStatusName : String
msg : String

updateTime : Date

idsCreated : List<String>

1   providerId, requestId

1..n
providerId

1
idsCreated

---

annotations
<AnnotationDocument>

🔑PK   _id : ObjectId

type : String

ownerId : String

dataSetId : String

1..n
dataSetid

---

dataSets <DataSetDocument>

🔑PK   _id : ObjectId

name : String
description : String
ownerId : String

dataBlocks : List<DataBlock> {
  beginTimeSeconds : Long
  beginTimeNanos : Long

1

0..n

DataBlock pvNames

```
<CommentAnnotationDocument>

        comment : String
```

The Data Platform MongoDB schema includes collections named "buckets", "requestStatus", "dataSets", and "annotations". Each is described in more detail below.

providers

The "provders" collection contains a document for each data provider registered with the Data Platform via the "registerProvider()" API method. Documents contain the following fields:

- "-id" - unique identifier for the provider
- "name" - name for provider as sent via registerProvider(), must be unique
- "attributeMap" - map of key/value metadata attributes describing the provider
- "lastUpdated" - contains time that the provider document was created or last updated (by calling registerProvider()).

buckets

The "buckets" collection manages bucketed time-series data for PVs, where each bucket contains a vector of PV measurements for a specified time range. The collection contains documents whose Java type is "BucketDocument".

The Ingestion Service manages the "buckets" collection. It creates a "BucketDocument" for each "DataColumn" (vector of PV measurements) in the "IngestionDataFrame" for a given "IngestDataRequest".

The domain of the Query Service methods is the "buckets" collection. A time-series data query specifies a list of PV names and time range and the result contains those buckets matching the query parameters. The metadata query returns results matching the bucket PV names.

Each bucket document includes the following fields:

- "_id" - unique identifier for the bucket
- "pvName" - name of the corresponding PV
- "providerId", "clientRequestId" - these fields are specified in the "IngestDataRequest" that created the bucket, used to find the corresponding requestStatus document with the disposition of the request
- "firstSeconds" / "firstNanos" and "lastSeconds" / "lastNanos" - time range for bucket with first and last time (with corresponding full Java Date fields "firstTime" and "lastTime" for convenience)
- "sampleCount" - specifies number of measurements contained in the "DataColumn" serialized to the bucket's dataColumnBytes field
- "samplePeriod" - specifies the sample period for the bucket in nanoseconds (zero if the "DataTimestamps" object for the bucket is a "TimestampsList", otherwise equal to the bucket's "SamplingClock.periodNanos")
- "dataTimestampsBytes" - contains serialized protobuf "DataTimestamps" object specifying the timestamps for the data vector values, using either a "SamplingClock" or "TimestampsList"
- "dataTimestampsCase" / "dataTimestampsType" - "DataTimestamps.value" enum case and name for the serialized object contained in dataTimestampsBytes (e.g., either a "SamplingClock" or "TimestampsList")
- "dataColumnBytes" - contains serialized protobuf "DataColumn" containing vector of values for bucket
- "dataTypeCase" / "dataType" - "DataValue.value" enum case and name for the first "DataValue" contained in the "DataColumn" for the bucket
- "attributeMap" - contains a list of key / value metadata pairs added to the bucket at ingestion

- "eventDescription" / "eventSeconds" / "eventNanos" - metadata description and start time for the event associated with this bucket at ingestion

requestStatus

The "requestStatus" collection contains documents whose Java type is "RequestStatusDocument". The Ingestion Service manages the "requestStatus" collection, creating a new document for each "IngestDataRequest" that it receives.

The service performs validation on each request, and responds with either a rejection or acknowledgment. The request is then handled asynchronously, with no further reporting back to the client making the request. The request status document indicates whether the request succeeded or failed, and provides further information for failures.

Given the asynchronous nature of the Ingestion Service, it is anticipated that a monitoring tool is needed to detect problems handling ingestion requests. Such a tool could use the "queryRequestStatus()" API method to query over the "requestStatus" collection.

Each request status document contains the following fields:

- "_id" - unique identifier for the request status
- "providerId" / "requestId" - these fields contain the corresponding values from the ingestion request and can be used by the client to match request status to the request
- "providerName" - contains the name for the corresponding providerId
- "requestStatusCase" / "requestStatusName" - IngestionRequestStatus enum case and name indicating the status for the request (e.g., rejected, error, success)
- "msg" - provides additional details for failed requests
- "updateTime" - specifies time status was updated
- "idsCreated" - contains a list of id strings for the buckets by a successful request
- "updateTime" - contains the time that the request status document was created

dataSets

The "dataSets" collection contains documents whose Java type is "DataSetDocument". As mentioned above, the annotation data model uses datasets to specify the relevant data for annotations.

The Annotation Service manages the "dataSets" collection. The API method "createDataSet()" creates a new document for each successful request. This collection is the domain for the method "queryDataSets()", which returns the documents in the collection matching the query criteria.

Each dataset document contains the following fields:

- "_id" - unique identifier for the dataset
- "name" - name of dataset
- "description" - a brief textual description of the dataset
- "ownerId" - owner of the dataset
- "dataBlocks" - contains a list of "DataBlock" objects, each of which contains a time range specified by "beginTimeSeconds" / "beginTimeNanos" / "endTimeSeconds" / "endTimeNanos" and a list of PV names

The time range and PV names specified in a "DataBlock" correspond to the domain of the "buckets" collection, and can therefore be used directly in the parameters for any of the Query Service time-series and metadata query methods.

annotations

The "annotations" collection contains documents whose Java type is "AnnotationDocument". This is a polymorphic hierarchy, with concrete classes extending that class for each type of annotation that is supported. Currently there is only one subtype, "CommentAnnotationDocument", though more will be added in an upcoming release.

The Annotation Service manages the "annotations" collection. The API method "createAnnotation()" creates a new document for each successful request. The handler for that method determines the appropriate concrete Java document class to create from the request parameters. This collection is the domain for the "queryAnnotations()" API method, which matches annotation documents against the search criteria and returns the matching documents.

The fields in the base document class "AnnotationDocument" include:

- "_id" - unique identifier for the annotation
- ○ "type" - specifies the discriminator string for the concrete document subtype, used by the MongoDB codec to map to the corresponding Java class
- ○ "ownerId" - string specifying the annotation owner
- ○ "dataSetId" - specifies the unique identifier string for the annotation's dataset